



INTERNATIONAL UNIVERSITY COLLEGE
LEUVEN

**DESIGN OF A WIRELESS SENSOR
NETWORKING TEST-BED**



Authors:

Bjorn DERAEVE
Roel STORMS

Supervisor:

Luc VANDEURZEN

April 2013

DESIGN OF A WIRELESS SENSOR NETWORKING TEST-BED

Bjorn Deraeve • Roel Storms

Written on behalf of Mr. Luc Vandeurzen

April 2013

Abstract

Audio applications are one of the main subfields of digital signal processing. More specific, in audio and music one can use frequency modulation synthesis in order to make creative and totally different sounding tones. A theoretical background and practical approach with an Analog Devices DSP-board is reported here.

In the first two chapters it is explained how sound works and what aspects of frequency modulation are actually responsible for making a particular tone sound differently. Next a practical approach to add envelopes and effects to the tone is given.

Chapter 5 serves as introduction to 6. In this chapter some theoretical background and electronic designs of CPU architectures are explained. This chapter serves as a backbone for the practical implementations explored in chapter 6. Readers with a basic knowledge on this matter can skip this chapter and immediately start with chapter 6. The first part of this chapter handles the Analog Devices ADSP-21369 DSP board's initialisation routines and the second part focusses on the practical implementation of the signal creation and modulation algorithms.

Finally chapter 7 deals with the communication between the DSP-board and the Labview interface for the FM Synthesizer.

As a conclusion chapter 8 closes this report with some critical reflections on the implementation and team activities.

Contents

Abstract	ii
List of Figures	vi
Abbreviations	ix
1 Introduction to Wireless Sensor Networks	1
1.1 Introduction	1
1.2 Wireless is everywhere	2
1.2.1 Definition of Wireless Sensor Network	2
1.2.2 Internet of Things	2
1.2.3 Third wave of wireless	2
1.3 Wireless Standards and Technology	2
1.3.1 Standards enable growth: ZigBee	2
1.3.2 Sensor nodes	3
1.3.3 Fundamental Challenges in WSNs	3
1.3.4 Power considerations	4
1.3.5 Network topologies	4
1.3.6 Security	4
1.4 Benefits of Wireless Measurements	4
1.5 WSNs: Innovative ways of interacting with the world	5
2 Why ZigBee	6
2.1 Introduction	6
2.1.1 Basic concepts of modulation synthesis	6
2.1.2 Oscillators and operators	7
2.1.3 Waveforms	7
2.2 Mathematical approach	8
2.3 Bandwidth considerations	10
2.4 Practical implementation	11
2.5 Summary	12
3 A WSN with Waspmotes: Theoretical aspects	13
3.1 Introduction	13
3.2 Hardware	14
3.2.1 Modular Architecture	14
3.2.2 Microcontroller and memory	14
3.2.3 Timers	16
3.2.3.1 Watchdog	16

3.2.3.2	Real Time Clock	16
3.3	Programming	16
3.4	Power considerations	16
3.4.1	Wasp mote power modes	16
3.4.1.1	Deep Sleep	17
3.4.1.2	Hibernate	17
3.4.2	Sampling sensors	18
3.4.3	Battery life estimation	18
3.4.3.1	XBee and Wasp mote start-up times	19
3.4.3.2	Battery life with standard program optimizations	19
3.4.3.3	Battery life with extra optimizations	21
4	A WSN with Wasp motes: Implementation aspects	23
4.1	Introduction	23
4.2	Libelium IDE and API	23
4.2.1	Wasp mote-IDE	23
4.2.2	Wasp mote-API	24
4.2.2.1	Original structure	24
4.2.2.2	Added functionality	25
4.2.3	Device start-up	26
4.2.3.1	Full initialization	26
4.2.3.2	Reduced initialization	26
4.3	Power savings	26
4.3.1	In the algorithm	26
4.3.2	Sleep vs. Hibernate	26
4.3.3	Variable sleep times	26
4.3.3.1	Calculate only the next time to sleep	27
4.3.3.2	Calculate all next times to sleep	27
4.3.3.3	Limit user control	28
4.3.4	Overdrive function	28
4.3.4.1	Influence of the overdrive effect on the frequency spectrum	29
4.3.4.2	Symmetrical clipping	29
4.3.4.3	Asymmetrical clipping	30
4.4	Equaliser	30
4.4.1	Design of the equaliser	30
4.4.2	Amplitude response of the FIR filters	30
4.4.3	Testing of the equalizer in MATLAB	30
4.4.4	Problems and solutions	32
5	Optimizing Wireless Sensor Networks	34
5.1	Introduction	34
5.1.1	What is digital signal processing?	34
5.1.2	How are DSP processors used?	34
5.2	Digitization	35
5.3	Digital Signal Processors	35
5.3.1	Numeric architecture	36
5.3.2	Memory architecture	37
5.3.3	Sequencer architecture	38
5.3.4	Input/Output architecture	38
5.4	Fixed-point vs. floating-point DSP	38
5.4.1	Advantages of floating-point digital signal processors	38
5.4.2	Considerations	38

6 Implementation: C	39
6.1 Introduction	39
6.2 Time management	39
6.2.1 Interrupts	39
6.2.2 Data handling	40
6.2.3 Real-time interface issues	40
6.3 Analog Devices Sharc DSP Board	41
6.3.1 Introduction	41
6.3.2 Initializing communication between the DSP and the audio jack	41
6.3.2.1 Digital Applications Interface (DAI)	41
6.3.2.2 Signal Routing Unit (SRC)	41
6.3.2.3 AD1835A	43
6.3.2.4 SPORT	45
6.4 The DSP Synthesizer program	46
6.4.1 Program execution	46
6.4.1.1 Not-realtime	46
6.4.1.2 Realtime: let's play	48
6.4.2 Key functions	49
7 Communication with LabVIEW User Interface using UART Port Controller	50
7.1 Introduction	50
7.2 UART data frame	50
7.3 UART external interface	51
7.4 UART configuration in DSP Processor for communication	52
7.5 Programming for communication between DSP Processor and LabVIEW User Interface	54
7.5.1 LabVIEW program for writing data	54
7.5.2 UART Interrupt	56
7.5.3 The FM synthesizer application	57
7.6 Conclusion	58
8 Critical reflections	59
8.1 Implementation problems	59
8.1.1 Communication with the Labview interface	59
8.1.2 Combining different program parts and general programming issues	59
8.1.3 Connecting the DSP board and use of the IDEs	60
8.1.4 Extreme programming	61
8.1.5 Latex	61
8.2 Team problems	61
8.3 Achieved activities	61
8.4 Future work	62
Bibliography	62
9 Bibliography	63
A XBee join and send time experiments	65
A.1 Distance relation test	65
A.2 First time reduction	67

List of Figures

1.1	Comparing the ZigBee standard with Cellular, Bluetooth and WiFi	3
1.2	Typical power consumption of a node	4
1.3	Typical power consumption of a node	4
2.1	Additive synthesis of a square wave	6
2.2	Basic waveforms	6
2.3	Square, triangular and sawtooth wave synthesis	7
2.4	Frequency spectrum of square wave	8
2.5	Frequency spectrum of triangular wave	8
2.6	Frequency spectrum of sawtooth wave	8
2.7	Phase function	9
2.8	The position of the side bands	9
2.9	The position of the side bands	10
2.10	Bessel functions for orders $n = 0,1,2, \dots$	11
2.11	The amplitude of the sidebands.	11
2.12	Functional design of the FM synthesizer program	12
3.1	Arduino XBee Shield	13
3.2	Libelium SquidBee	13
3.3	Wasp mote V1.1 with Bluetooth expansion module	14
3.4	Wasp mote PRO	14
3.5	Main Wasp mote components	15
3.6	Wasp mote block diagram	15

3.7	Wasp mote going from ON to Deep Sleep	17
3.8	Wasp mote going from ON to Hibernate	18
3.9	Battery life in High performance mode	20
3.10	Energy usage in High performance mode	20
3.11	Battery life High Performance vs. Power Saver	21
3.12	Energy usage in Power Saver mode	22
4.1	Reduced dependency graph for WaspXBeeZB	25
4.2	Reduced dependency of core libraries	25
4.3	Effect of overdrive in time domain	28
4.4	Overdrive threshold 1	29
4.5	Overdrive threshold 0.5	29
4.6	FFT of a basic signal	29
4.7	Symmetric clipping	30
4.8	Asymmetric clipping	30
4.9	Architecture of the equaliser	31
4.10	Magnitude response lowpass filter	31
4.11	Magnitude response bandpass filter	31
4.12	Magnitude response highpass filter	32
4.13	$\text{Gain}_{LP} = 0$, $\text{Gain}_{BP} = 1$, $\text{Gain}_{HP} = 1$	32
4.14	$\text{Gain}_{LP} = 1$, $\text{Gain}_{BP} = 0$, $\text{Gain}_{HP} = 1$	33
4.15	$\text{Gain}_{LP} = 1$, $\text{Gain}_{BP} = 1$, $\text{Gain}_{HP} = 0$	33
5.1	DSP Principle	35
5.2	Signal sampling representation	35
5.3	Structere of a DSP	36
5.4	8-bit Barrel Shifter	36
5.5	DSP architecture	37
5.6	Harvard vs. Von Neuman	37
6.1	Serial communication interface	40

6.2	System Architecture Block Diagram	42
6.3	Peripherals Architecture Block Diagram	42
6.4	Pin Buffer	43
6.5	AD1835A on the ADSP-21369 DSP Board	44
6.6	Peripherals Architecture Block Diagram	44
6.7	Flow chart of the DSP Synthesizer program	47
7.1	A typical UART data frame	51
7.2	System Architecture Block Diagram	51
7.3	Schematic for connection of UART port	52
7.4	UART Functional Block Diagram	53
7.5	Write/Read a string to and from a COM port	54
7.6	Formatting a string before writing to DSP	55
7.7	A part of writing parameters to Operator 1 program	55
7.8	PICR2 Register	56
7.9	VI Hierarchy of LabVIEW User Interface	57
7.10	Front Panel of the application	58
7.11	The algorithm for three operators	58
A.1	XBee join and send times: distance relation	66
A.2	XBee join and send times: first time reduction	67

Abbreviations

API	Application Programming Interface
AVR	Alf Vegard RISC
CPU	Central Processing Unit
DARPA	Defense Advanced Research Projects Agency
DRAM	Dynamic Random Access Memory
DSN	Distributed Sensor Network
EEPROM	Electrically Erasable Programmable Read Only Memory
ENIAC	Electronic Numerical Integrator And Computer
IDE	Integrated Development Environment
IOT	Internet Of Things.
OEM	Original Equipment Manufacturer
PAN	Personal Area Network
PCB	Printed Circuit Board
RFID	Radio Frequency IDentification
RISC	Reduced Instruction Set Computing
RTC	Real Time Clock
SD	Secure Digital
SRAM	Static Random Access Memory
UPC	Universal Product Code
WSN	Wireless Sensor Network

Chapter 1

Introduction to Wireless Sensor Networks

What if: Ubiquitous Computing & Networking + Sensing & Control ?

1.1 Introduction

The ENIAC, the first electronic computer designed by the American scientists J. Presper Eckert and John W. Mauchly, Turing-complete, pioneering in 1946 and designed to calculate artillery firing tables. A wireless sensor network node, over 50 years later and 10 million times less the size of the ENIAC, still Turing-complete, and still military roots.

The origins of the research to Wireless Sensor Networks (WSNs) can be traced back to DARPA¹, which sponsored a workshop at Carnegie Mellon University in 1978, identifying the technology components for a DSN. But at the time the technology was not quite ready. Sensors could take up the size of a shoe box and up, limiting the number of potential applications. The earliest DSNs were also not very tightly associated with wireless communication.

In 1998 a new wave of research in WSNs started. Again DARPA acted as pioneer by launching the initiative research program 'SensIT', which added new capabilities to the current sensor network such as ad hoc networking, dynamic querying and tasking, reprogramming and multi-tasking. At the same time IEEE noticed the high capabilities and low expenses of WSNs and defined the IEEE 802.15.4 standard for low power consumption, low data rate wireless PANs. for 868MHz, 915MHz and 2.4GHz radios. Finally, in 2002, the ZigBee Alliance was established and published the ZigBee standard, based on IEEE 802.15.4. The standard adds a suite of high level communication protocols for WSNs such as device coordination, network topologies and interoperability with other wireless products.

Currently, WSNs are seen as one of the most prospective technologies of the 21st century (21 Ideas for the 21st Century, 1999). China for example, has involved WSNs in their national strategic research programmes (Program 973). The project follows an application-driven methodology and aims to issues identified with the real-world critical problems facing Chinese society. Over the years the program has funded areas such as agriculture, health, resources, energy, population and materials and brought significant benefits to China's sustainable economic and social development.

¹Defense Advanced Research Projects Agency (DARPA), is an American agency responsible for developing military technologies. DARPA has fund the research in many technologies which have had major effect on the world, including ARPANET, the first wide-area packet switching network (ancestor of the Internet and Douglas Engelbart's precursors to the GUI.

1.2 Wireless is everywhere

1.2.1 Definition of Wireless Sensor Network

A *Wireless Sensor Network (WSN)* consists of spatially distributed autonomous sensors connected via a (wireless) communications infrastructure to cooperatively monitor, record and store physical or environmental conditions, such as temperature, sound, vibration, pressure, motion or pollutants (Sebastian Büttrich)

1.2.2 Internet of Things

The Internet of Things is a term that was first used by Kevin Ashton² in 1999 and answers the question introducing this chapter. It refers to uniquely identifiable objects (things) and their virtual representations in an Internet-like structure. It is a vision of a network of Internet-enabled objects, together with web services which interact with these objects. If all objects in the world would be equipped with minuscule identifying devices this could transform our daily lives. By embedding computational capabilities in all kinds of objects, including living beings, it will be possible to provide qualitative and quantitative shift in several sectors: logistics, domotics, healthcare, entertainment, and so on.

WSNs can provide a virtual layer where the information about the physical world can be accessed by any computational system. As a result, WSNs are one of the most important elements for realizing the vision of the Internet of Things paradigm. On May 2nd, 2012, Libelium³ published a list of 50 cutting edge *Internet of Things* applications . According to Libelium, by 2020, more than 50 billion devices will be connected to the Internet.

1.2.3 Third wave of wireless

The IoT can also be considered as the third wireless wave. Following cellular technology and WiFi, today wireless technology also includes the world of sense and control technology to bridge the gap between the virtual electronic world and our human physical world. This wave is not focussed on enhancing human communication. Instead it allows sensors to interact with actuators, creating a more dynamic world and avoiding costly human labour. Finally, as the title of this section suggests, WSNs offer the ability to observe the previously unobservable.

1.3 Wireless Standards and Technology

1.3.1 Standards enable growth: ZigBee

Not only do standards allow devices from different vendors to interoperate, they also provide OEMs and integrators the flexibility of second sourcing. The ZigBee Alliance is an independent standardization organization which is driven by a large group of OEM companies and has definitely had a large impact on the rapid development of WSNs.

Figure 1.1 indicates the most critical properties of the ZigBee standard. Some rules-of-thumb are:

²Kevin Ashton co-founded the Auto-ID Center at the Massachusetts Institute of Technology (MIT). The Center is a research group in networked RFID and newly emerging sensing technologies. The main goal was the development of the Electronic Product Code (EPC), a global RFID-based item identification system intended to replace the UPC bar code

³The manufacturer of the Hardware we used. See chapter 3 for more information.

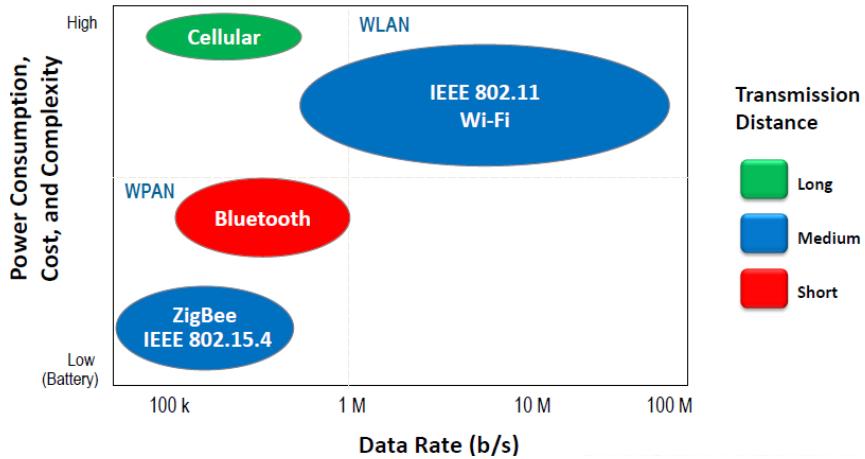


FIGURE 1.1: Comparing the ZigBee standard with Cellular, Bluetooth and WiFi

- The higher the frequency, the higher the data rate and power consumption
- The lower the frequency, the further the reach
- All radio waves show strong absorption in water and metal

Table 1.1 shows the typical power consumption, throughput, range and application examples of each technology:

	Battery Life	Data Rate	Range	Application Examples
ZigBee	1-4 years	20 to 250Kbps	100 m	Wireless Sensor Networks
Bluetooth	1-2 weeks	1 to 3 Mbps	10 m	Wireless Headset
IEEE 802.11g	1-2 days	6 to 54Mbps	30 m	Wireless Internet Connection

TABLE 1.1: Comparing the ZigBee standard with Cellular, Bluetooth and WiFi

1.3.2 Sensor nodes

Sense Compute Store Communicate

1.3.3 Fundamental Challenges in WSNs

The biggest challenge WSNs encounter is without a doubt power consumption. Power! Power! Power! The next section will briefly introduce the general concepts and in section ?? the power consumption of the WSN we developed will be analyzed thoroughly. Other challenges often fall back to this limited amount of available energy, security for example.

There are dozens of other basic challenges for WSNs. Unattended operation and environmental influence makes a mote prone to failure. Mobility can cause topology changes. WSNs can use over more than 100 nodes, this leads to scalability issues and synchronization issues. There must also be a form of synchronization with sleeping nodes. Network responsiveness and robustness (this will be discussed more in section 3.4. Not to forget making sense out of sensors).

1.3.4 Power considerations

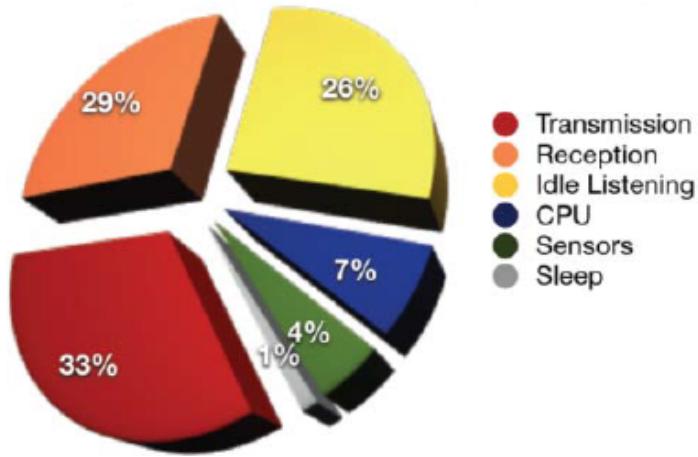


FIGURE 1.2: Typical power consumption of a node

Node Power Consumption

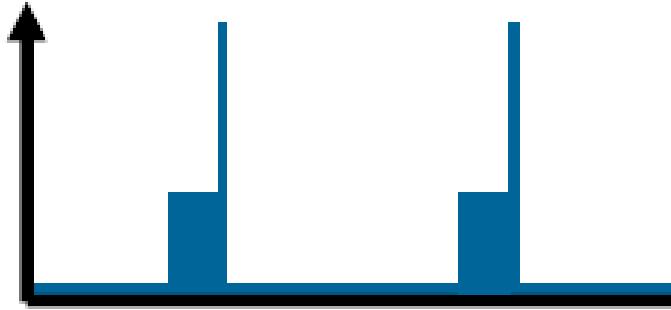


FIGURE 1.3: Typical power consumption of a node

1.3.5 Network topologies

1.3.6 Security

1.4 Benefits of Wireless Measurements

Wireless Sensor Networks provide benefits in roughly three categories:

1. Reduce costs
 - Installation costs and time
 - Maintenance costs
2. Increase efficiency
 - Optimize measurement processes

- Access data everywhere and every time
 - Decrease downtime
3. Monitor anywhere
 - Overcoming power or infrastructure limitations
 - Solve new and previously challenging applications

1.5 WSNs: Innovative ways of interacting with the world

pdf intro to sensornets

Chapter 2

Why ZigBee

Written by Bjorn Deraeve

2.1 Introduction

Thanks to new digital technology there was room for a new synthesis technique, Frequency Modulation Synthesis. The advantage of FM is that a very large number of different sounds can be produced with few elemental units (oscillators). In other words the timbral space is very large. The one requirement is the use of optimized computers, which is discussed in chapter 5.

2.1.1 Basic concepts of modulation synthesis

It is important to know that in the synthesizer only one wave form, a sine wave, is stored. The trick of modulation synthesis is to remodel this wave until a desired sound is found.

An easy example is the square wave. Using additive synthesis a square wave can be built out of an infinite amount of sine functions. This example of additive synthesis also explains the Fourier transformation theory and is shown in figure 2.1.

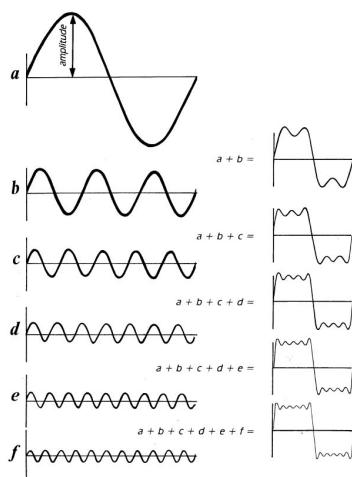


FIGURE 2.1: Additive synthesis of a square wave

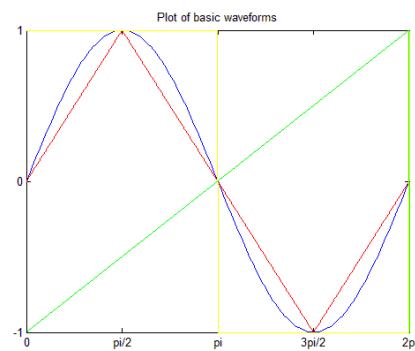


FIGURE 2.2: Basic waveforms

2.1.2 Oscillators and operators

The basic building blocks of FM synthesizers are operators. They are equivalent to oscillators in analogue synthesizers and produce the waveforms the user desires. Where an analogue oscillator produces a changing voltage the digital operator creates discrete samples according to a sine pattern. The operators invite the user to be creative in creating their own unique sound with the aid of different kinds of modulation techniques and waveforms. Knowing that most interesting sounds have many components, with the use of only four operators and modulation techniques one can generate sounds with much more than four partials.

2.1.3 Waveforms

To produce creative sounds it can be appropriate to combine sine waves with other waveforms. Synthesizers are typically able to produce sines, triangles, sawtooths and square waves. A square wave contains only odd-integer harmonic frequencies. To be an ideal wave, the signal should change from high to low instantaneously. However this is impossible to achieve in real-world systems because of the limited bandwidth, though good approximations can be made with additive synthesis. The Fourier series of the square wave is given in equation 2.1

$$x(t) = \frac{4}{\pi} \sum_{n=1,3,5,\dots}^{\infty} \frac{\sin(2\pi \cdot n \cdot f \cdot t)}{n} \quad (2.1)$$

$$x(t) = \frac{4}{\pi} \left(\sin(\omega t) + \frac{1}{3}\sin(3\omega t) + \frac{1}{5}\sin(10\omega t) + \dots \right) \quad (2.2)$$

Triangle waves contain just like square waves only odd harmonics of the pitch. The difference is that the higher harmonics lose their energy faster. Equation 2.3 gives the Fourier series to approximate a triangle wave using additional synthesis.

$$x(t) = \frac{8}{\pi^2} \sum_{n=1,3,5,\dots}^{\infty} (-1)^{(n-1)/2} \frac{\sin(2\pi \cdot n \cdot f \cdot t)}{n^2} \quad (2.3)$$

$$x(t) = \frac{8}{\pi^2} \left(\sin(\omega t) - \frac{1}{9}\sin(3\omega t) + \frac{1}{25}\sin(5\omega t) - \dots \right) \quad (2.4)$$

Finally sawtooth waves contain both even and odd harmonics of the fundamental frequency. The sound of it is even more raw. The equation for additive synthesis is given by equation 2.5.

$$x(t) = \frac{2}{\pi} \sum_{n=1}^{\infty} (-1)^{n+1} \frac{\sin(2\pi \cdot n \cdot f \cdot t)}{n} \quad (2.5)$$

$$x(t) = \frac{2}{\pi} \left(\sin(\omega t) - \frac{1}{2}\sin(2\omega t) + \frac{1}{3}\sin(3\omega t) - \dots \right) \quad (2.6)$$

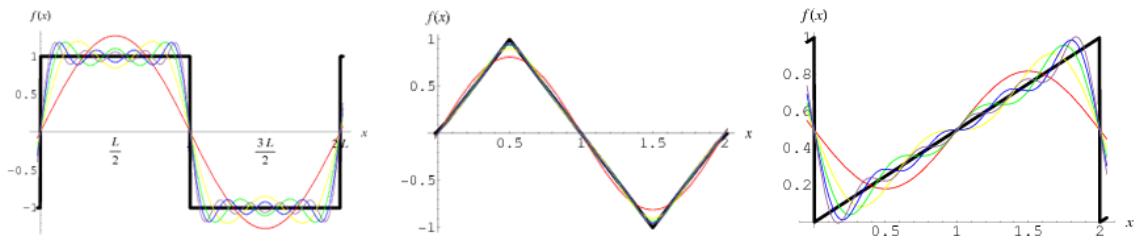


FIGURE 2.3: Square, triangular and sawtooth wave synthesis

Figures 2.4, 2.5 and 2.6 show the corresponding frequency responses of the square, triangle and sawtooth wave and confirm equations 2.1, 2.3 and 2.5.

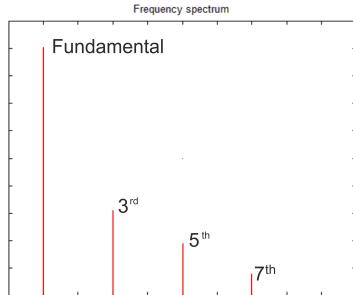


FIGURE 2.4: Frequency spectrum of square wave

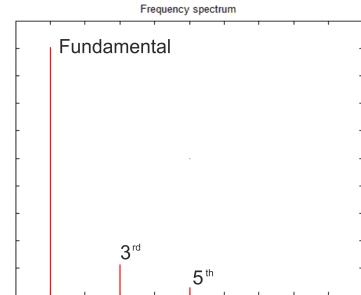


FIGURE 2.5: Frequency spectrum of triangular wave

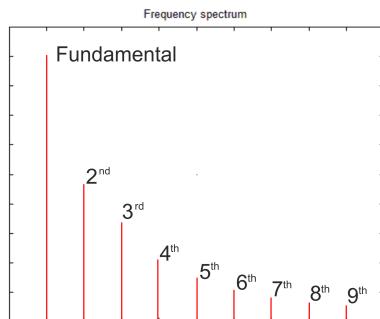


FIGURE 2.6: Frequency spectrum of sawtooth wave

2.2 Mathematical approach

The most basic form of frequency modulation is vibrato, in fact, FM is simply very fast vibrato. This possibility to create completely new sounds has been discovered by John Chowning at Stanford University in the mid-'60s. He found that when the frequency of the modulator increases beyond a certain point, the vibrato effect is replaced by a complex new tone.

We can create a vibrato sound by letting operator 2 cause a small change in frequency of operator 1. The rate and amount of frequency change is determined by operator 2. When the waveform of operator 2 increases the tone becomes brighter. This means that more harmonics are appearing in the frequency spectrum. Referring to the introduction section of this chapter, we remark again that FM is able to produce a wide variety of interesting spectra with only two basic oscillators.

Equations 2.7 and 2.8 are two basic sines. The instantaneous amplitude of the waves at any given point in time is called 'A' and are related to their gains 'a' (maximum amplitude of their cycle).

$$A_c(t) = a_c \sin(2\pi f_c t) \quad (2.7)$$

$$A_m(t) = a_m \sin(2\pi f_m t) \quad (2.8)$$

Equations 2.9 and 2.10 define the frequency modulation of the carrier wave $A_c(t)$ by the modulation wave $A_m(t)$.

$$A_c(t) = a_c \sin((2\pi f_c + A_m) t) \quad (2.9)$$

$$A_c(t) = a_c \sin((2\pi f_c + a_m \sin(2\pi f_m t)) t) \quad (2.10)$$

Looking at equation 2.10 we see the modulated signal consists of a constant amplitude a_c and a time varying argument to the sine function, or phase function $\phi(t)$. This phase function is the addition of a constant ramp with slope $2\pi f_c$ and a sinusoidal variation, see picture 2.7. The

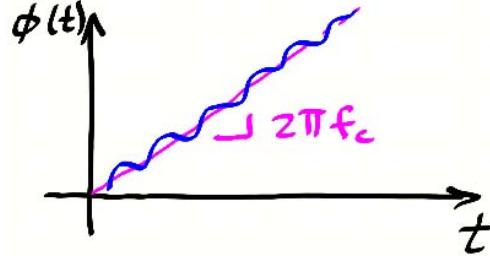


FIGURE 2.7: Phase function

phase function has two parameters that influence the nature of the modulated signal: a_m and f_m . The degree to which this variation is added to the constant ramp is given by the modulation index a_m or also referred to as I. So the modulation index describes the amount of frequency deviation. f_m is the modulation frequency and defines the rate at which the frequency deviation is to occur.

Now how does this influence the sound of a frequency modulated signal? First one must see that the side bands produced by FM lie around the carrier frequency, being plus or minus an integer multiple of the *modulator frequency* (the 1st parameter of $\phi(t)$). This is expressed by equation 2.11 and visible in figure 2.8.

$$f_{sb} = 2\pi f_c \pm n \cdot 2\pi f_m \quad (2.11)$$

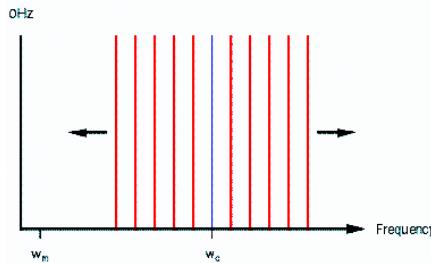


FIGURE 2.8: The position of the side bands

The other parameter of the phase function is the *amplitude of the modulator signal* and has a direct relation to the modulation index I. The modulation index is defined as the ratio of the change in carrier frequency to the modulator frequency:

$$I = \frac{\Delta\omega_c}{\omega_m} \quad (2.12)$$

Obviously the change in carrier frequency is determined by the amplitude of the modulator signal a_m . For any given modulator frequency, it is the modulation index (and thus the amplitude of the modulator) that defines the energy of each of the components in the spectrum. Also the height of the carrier frequency can be influenced by the value of I.

In fact the amplitude of the sidebands is controlled by a Bessel function $J_k(a)$. Equation 2.10 can be rewritten as an infinite summation given by equation 2.13, in which θ corresponds to $2\pi f_c$, a corresponds to a_m (strictly spoken I) and β corresponds to $2\pi f_m$:

$$\sin(\theta + a \sin \beta) = J_0(a) \sin \theta + \sum_{k=1}^{\infty} J_k(a) [\sin(\theta + k\beta) + (-1)^k \sin(\theta - k\beta)] \quad (2.13)$$

In this equation $\sin(\theta)$ corresponds to the component based on the carrier wave and the series of $\sin(\theta + k\beta)$ are for the side band components. Note that the amplitude of all the components is controlled by the Bessel functions.

First we will show exactly how the frequency of the modulator signal defines the position of the side band components: in equation 2.13 we see "sin(θ)", this leads to the frequency component of the carrier f_c , see picture 2.9. If k equals one, the sine functions between brackets lead to frequency components of the modulator, $f_c + f_m$ and $f_c - f_m$. If k equals two components $f_c + 2f_m$ and $f_c + 2f_m$ are formed. And so on. The effect of the negative components will be ignored from now on¹. Now we will have a look at how the Bessel functions control the amplitude of the

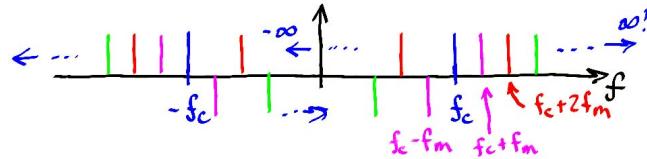


FIGURE 2.9: The position of the side bands

side bands. Equation 2.14 defines a Bessel function with order α .

$$x^2 \frac{d^2y}{dx^2} + x \frac{dy}{dx} + (x^2 - \alpha^2)y = 0 \quad (2.14)$$

As defined in equation 2.13, the order of the Bessel function corresponds to the sideband number. So figure 2.10 shows 6 Bessel functions as a function of modulation index for a specific sideband number.

Note that, if the modulation index is zero (no modulation), the amplitude of the carrier frequency is equal to 1 (black curve). As we look at sidebands that are farther away from the carrier frequency we see that it takes longer for the amplitude of that sideband to become non-zero (red, blue,... curves). So one Bessel curve shows that the amplitude of the n^{th} sideband depends on the modulation index. Or, for a given modulation index, the amplitude of the n^{th} sideband is given by a Bessel function of order n .

When the modulation index of the signals in figure 2.8 is increased the spectrum will look like figure 2.11:

Appendix A contains a simple labview program to display Bessel functions of different orders and more examples that summarize this section.

2.3 Bandwidth considerations

The bandwidth of the signal can be defined as the range of frequencies occupied by the signal. Although the sum series of sidebands is theoretically infinite, the modulation index will make sure that sidebands of higher frequency are very small and negligible.

Equation 2.15 gives an approximation of the bandwidth of a FM signal.

$$B = 2\pi f_m(1 + I) \quad (2.15)$$

¹Since we deal with an infinite summation, the negative components will kind of mix into the positive components. This can be thought of as a kind of aliasing phenomenon.

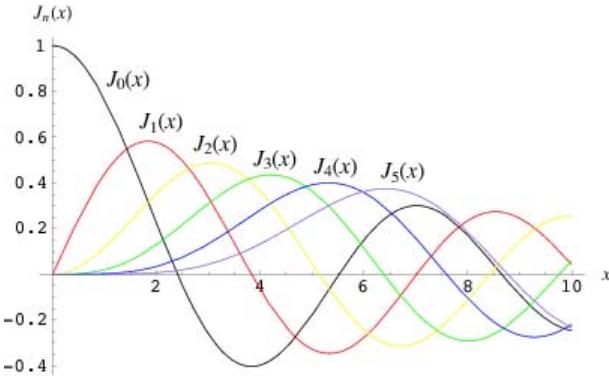


FIGURE 2.10: Bessel functions for orders $n = 0, 1, 2, \dots$. The order defines the curve, x-axis defines the modulation index and the y-axis returns the amplitude.

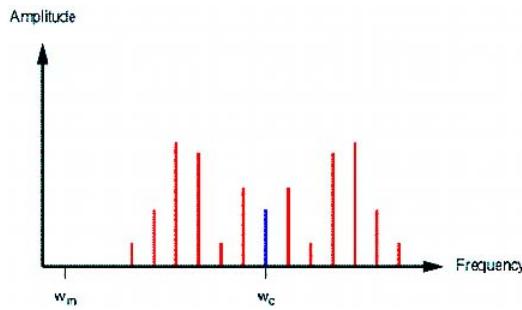


FIGURE 2.11: The amplitude of the sidebands.

For example: the bandwidth of a signal with a 300Hz modulator with $I = 5$ will be $2 \times 300\text{Hz} \times (1+ 5) = 3600\text{Hz}$. This show again that with FM it is easy to create complex signals with a much higher bandwidth.

2.4 Practical implementation

Figure 2.12 shows the practical implementation of the FM theory in our synthesizer program. In this figure the first modulation mode is shown where operator 2 modulates operator 1 according to formula 2.10. The user can choose from this list of modulation types:

- 2 \rightarrow 1
- 3 \rightarrow 2 \rightarrow 1
- 4 \rightarrow 3 \rightarrow 2 \rightarrow 1
- 4 \rightarrow 3 + 2 \rightarrow 1

For more information on the implementation of the FM synthesizer program please read chapters 5, 7 and 7.

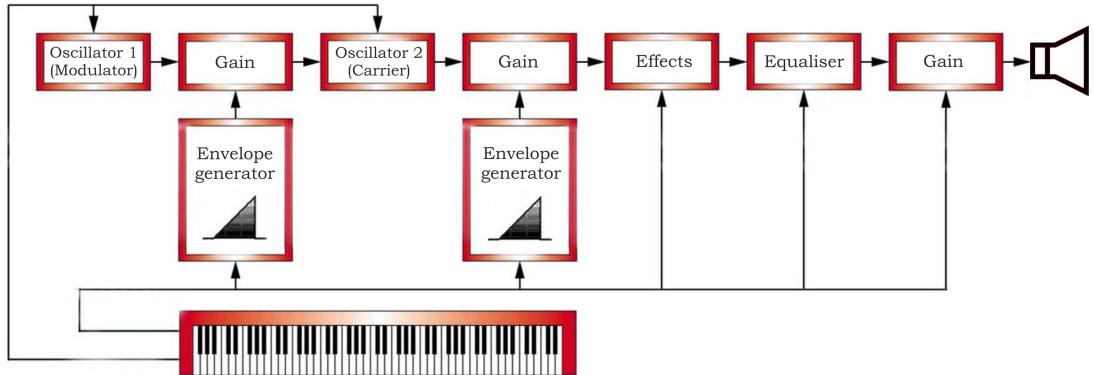


FIGURE 2.12: Functional design of the FM synthesizer program
 (First modulation type from the list above)
 (Compare this figure with the program flow chart 6.7 discussed in chapter 6.)

2.5 Summary

We've showed that frequency modulation is a very powerful method for music synthesis. It is possible to generate sounds that are not obtainable by any other modulation technique.

The carrier frequency sets the center of the sideband cluster, the modulator frequency only sets the sideband positions.

It is the modulation index that determines the amplitude of the sidebands and the number of significant sidebands, doing so shaping the spectrum.

Chapter 3

A WSN with Waspmotes: Theoretical aspects

3.1 Introduction

Wasp mote is more than just another piece of hardware. In fact it is an open source platform for wireless sensors, specially focusing on low consumption and autonomy. Wasp motes promise to offer a variable lifetime between 1 and 5 years, depending on the duty cycle and the used radio. But it didn't just start with Wasp mote and it will definitely not end with it. In 2007 developers from Libelium collaborated with the Arduino Team creating the first open hardware shield for Arduino, the "Arduino XBee Shield". The shield allowed an Arduino board to communicate wirelessly via ZigBee. Libelium used the shield to develop their first sensor device, the SquidBee, intended for creating sensor networks. Although the SquidBee is self-powered and implements wireless communications, it is more sensor device than wireless sensor device. The 3.3V - 5V regulator could not be turned off, as a result there is a constant consumption of 50mA discharging the battery within hours. The SquidBee was created for teaching and educational purposes only. Since the platform was not radio certified the motes could not be deployed in real scenarios like cities, factories or even houses, so it did not fit Libelium's corporate customer requirements at all. However, the tone of an open hardware and source wireless sensor device was definitely set.

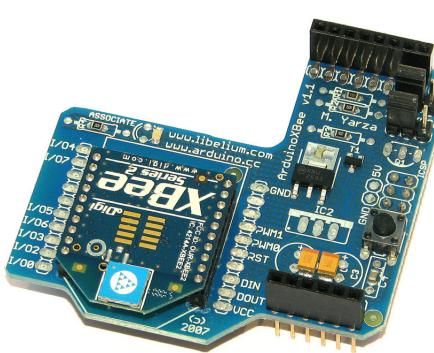


FIGURE 3.1: Arduino XBee Shield



FIGURE 3.2: Libelium SquidBee

In 2009 the Wasp mote was born, meeting all the above requirements: low consumption and meeting three radio certification requirements (CE for Europe, FCC for the US and IC for Canada). In addition the Wasp mote was built with a complete modular philosophy. The idea behind this design is to integrate only the needed modules in each device, optimizing costs. This is why all modules are connected to the Wasp mote via sockets.

Since its introduction, more than 2000 developers have been using Waspmote (v1.1) and the platform has received many suggestions and possible improvements. Libelium carefully listened to all these proposals and decided to bring out a new version with the name of Waspmote PRO (v1.2) in February 2013. This new version comes with upgraded hardware and an improved API, which is unfortunately not compatible with the older API. The most important improvements of the hardware is that the code can be uploaded much quicker and the XBee radio must no longer be removed, which is a huge improvement. There are also no more jumpers and there is no need of a coin battery. Regarding the API, Libelium claims it is much more robust and easier to use than the previous one and they also improved their programming guides.

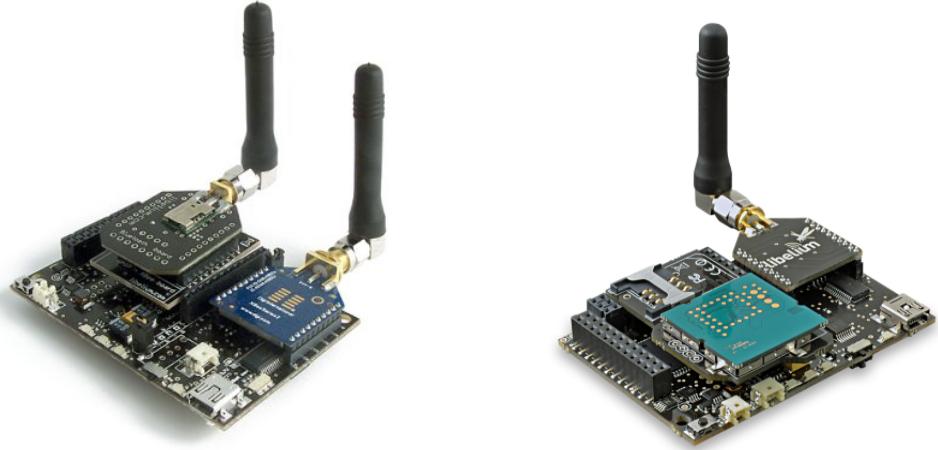


FIGURE 3.3: Waspmote V1.1 with
Bluetooth expansion module

FIGURE 3.4: Waspmote PRO

3.2 Hardware

3.2.1 Modular Architecture

As mentioned in this chapter's introduction, Waspmote is based on a modular architecture, doing so optimizing costs and able to change according to the specific user's requirements. The available modules are split up into five categories:

- ZigBee
- GSM - 3G/GPRS
- GPS
- Sensor Boards
- Storage

Figure 3.5 indicates the Waspmote main components.

3.2.2 Microcontroller and memory

Just like on any other PCB, the CPU is the heart of the module. Waspmote integrates an 8-bit ATmega 1281 microcontroller with 128KB programmable flash, 8KB SRAM runtime memory

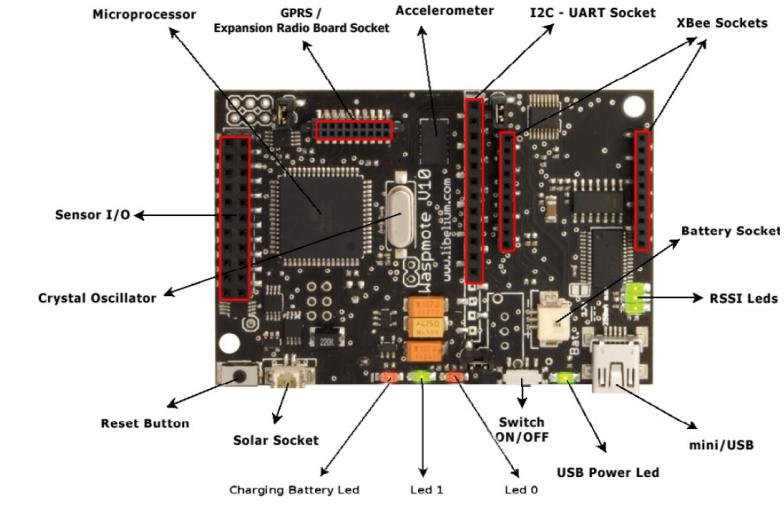


FIGURE 3.5: Main Wasp mote components

and 4KB EEPROM memory. Since SRAM is built with cleverly combined MOSFETs it must not be periodically refreshed, but it is still volatile memory. The main advantage compared to DRAM is that, when moderately clocked like in the Wasp mote, it consumes very little power. Because of the modular design of the Wasp mote the block diagram is very simple. This is shown in figure 3.6.

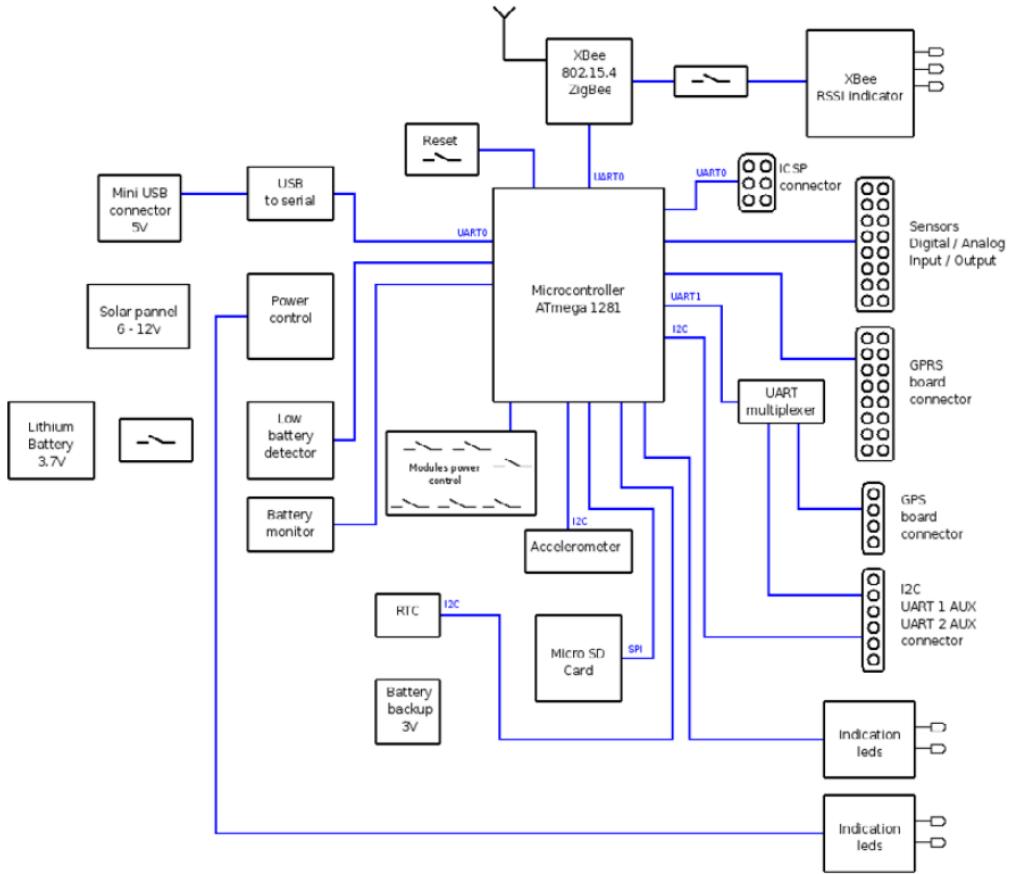


FIGURE 3.6: Wasp mote block diagram

3.2.3 Timers

The Waspmove's system clock is an 8MHz quartz oscillator. This means that every 125ns a machine language instruction is executed by the microcontroller. Keep in mind that one C++ instruction consists of several instructions in machine language. To generate interrupts the Waspmove has an internal watchdog and a Real Time Clock (RTC).

3.2.3.1 Watchdog

The Watchdog is integrated on the Atmega 1281 and counts the clock cycles generated by a 128KHz oscillator. When the WDT counter reaches a set value it generates an interruption signal. The WDT is used to awake the microcontroller from *Sleep* mode, because of its high precision. Thus, *Sleep* mode allows small intervals, going from 16 milliseconds to a maximum of 8 seconds.

3.2.3.2 Real Time Clock

To store an absolute time base the RTC can be used. Alarms programmed in the RTC specify days, hours, minutes and seconds. For the RTC the waspmote uses a Maxim DS3231SN 32.768Hz oscillator. Because this clock has an internal compensation mechanism for variations caused by temperature changes this is one of the most accurate clocks on the market. This clock is used to wake the Waspmove from the higher energy saving modes *Deep Sleep* and *Hibernate*, with intervals from 8 seconds to even days. It is important to notice that in *Hibernate*, the RTC is no longer powerd through the main battery but through the auxiliary (button) battery. So when problems occur when using hibernate probably it is recommended to measure the button battery's voltage and possibly must be replaced.

3.3 Programming

To develop on the Waspmove, Libelium offers an API and IDE. For more information on the API, please see chapter 4. Waspmove uses the same IDE (compiler and core libraries) as Arduino, so as long as things like pin layout and I/O schemes are adjusted, code should be compatible in both platforms. So Waspmove and Arduino are pretty much the same, however don't forget that Waspmove has Radio Certifications and Arduino doesn't.

When the Waspmove is started, the microcontroller will execute the bootloader and start loading the compiled program from FLASH into the SRAM working memory.

The code is divided into two basic parts: **setup** and **loop**, each with sequential behaviour. When the Waspmove is switched on or reset, the code starts at the setup function and then enters the loop function. Because the second part forms an infinite loop a common technique to save energy is to block the program until some interruption is detected.

Since in *Hibernate* mode the Waspmove is completely disconnected from the main battery also the program is interrupted. This means the SRAM has lost all variable values and at wake up the code restarts at the **setup** function. To store values during hibernate cycles it is necessary to write them to EEPROM or the SD card.

3.4 Power considerations

3.4.1 Waspmove power modes

The libelium Waspmove has 4 operational modes: ON, Sleep, Deep Sleep and Hibernate. They differ from which type of interruptions they can be woken up and duration interval. For our

application we want sleep intervals of 30 seconds and more, so only Deep Sleep and Hibernate mode are of interest. Table 3.1 summarizes the Waspmotes operational modes.

Mode	Consumption	CPU	Cycle	Accepted Interruptions
ON	9mA	ON	-	Synch and Asynch
Sleep	62µA	ON	31ms - 8s	Synch (WDT) and Asynch
Deep Sleep	62µA	ON	8s - min/hours/days	Synch (RTC) and Asynch
Hibernate	0.7µA	OFF	8s - min/hours/days	Synch (RTC)

TABLE 3.1: Operational modes of Libelium Waspmove V1.1

3.4.1.1 Deep Sleep

In deep sleep mode the main program is paused and the CPU passes to a latent state. Triggers are as well synchronous interruptions (RTC) as asynchronous interruptions. Examples of asynchronous interruptions are low battery level or a sensor that reaches a certain trigger value.

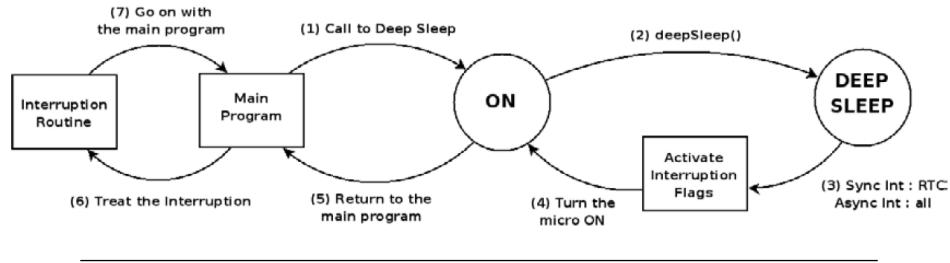


FIGURE 3.7: Waspmove going from ON to Deep Sleep

In figure 3.7 the process from going to operational mode ON to Deep Sleep is shown. The main advantage of this mode is that the program is only paused, so the program stack and thus all variable values keep their values. When the Waspmove is turned back on it simply executes the next instruction.

3.4.1.2 Hibernate

Hibernate mode consumes roughly 100 times less energy than Deep Sleep. This is made possible by disconnecting all the Waspmove's modules, including the microcontroller. The RTC gets his power through the auxiliary battery. So if hibernate mode stops working it is probably necessary to replace the Waspmove's button battery. Figure 3.8 demonstrates the process from ON to hibernate.

This means the CPU is also switched off and does not remember any values from variables. When waking up the Waspmove reinitializes, the microprocessor is reset and the program restarts from the beginning. Both **setup** and **loop** routines are executed as if the main switch would be activated. By placing the **ifHibernate()** function in setup the program can determine if it came from a hardware reset or from a hibernate reset. To be able to wake up from hibernate mode the hibernate jumper must be removed correctly. See section for remarks on this issue. Because not all Libelium's API functions regarding hibernate in combination with the different alarm modes work, it is advised to use the functions provided in **WaspXBeeZBNode.h**.

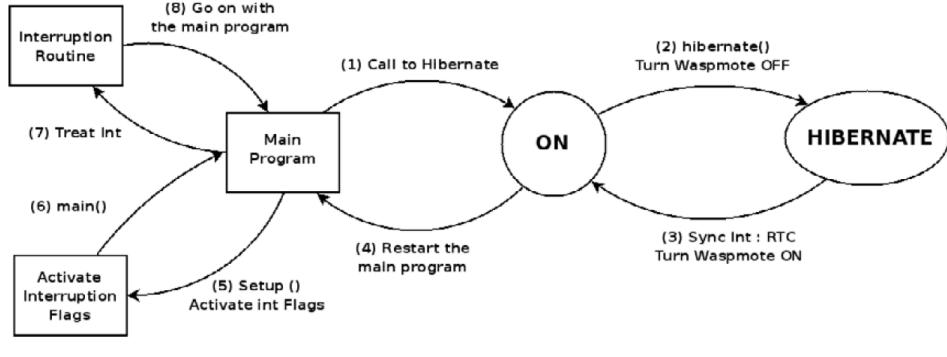


FIGURE 3.8: Waspmote going from ON to Hibernate

3.4.2 Sampling sensors

To measure the sensors, originally we took 10 samples with 100 milliseconds recommended delay between the measurements and calculated the average. Since we want to make the energy consumption as low as possible we now do the iterations without delay. Appendix ?? contains 60 test samples per sensor and indicates that there is no significant difference on the average by removing this delay. Except for CO₂ measurements, removing the delay saves about 1000 milliseconds per measured sensor.

Because the first sample often shows a slight deviation, the program takes 11 samples but bases the average on the last 10 values.

3.4.3 Battery life estimation

In order to be able to give recommended sensor measuring intervals this section will analyse the estimated battery life of the Waspmotes. Table 3.5 enumerates the most common components typical consumption. The batteries included with our Waspmotes are rechargeable Lithium-ion

Action	Average Current
XBee, sending,	105mA
CO ₂	50mA
XBee, ON	45mA
Waspmote, ON	9mA
Pressure	7mA
Humidity	380μA
Waspmote, sleep	62μA
Temperature	6μA
Waspmote, hibernate	0,007μA

TABLE 3.2: Operational modes of Libelium Waspmote V1.1

batteries with a capacity of 6600mAH. The Waspmote that will be deployed on the roof of Group T, campus Vesalius will also have a 12V solar panel with a charging current up to 280 mA to extend its battery life. The other batteries can be charged manually or by USB (5V, 100mA). Lithium-ion have a self-discharge rate of typically 1 to 2 percent per month and since the used batteries are new we expect a high battery efficiency.

3.4.3.1 XBee and Waspmove start-up times

For the XBee node to join an existing network there are two power related possibilities. Either the Waspmove has been turned on already sufficiently long and the XBee had more than enough time to join the network, or either the XBee wasn't joined yet and the program needs to wait on this. From the experiments done at our apartment we came to following conclusions:

1. It takes about 2.5 seconds to join a network after powered on.
2. If the XBee is joined, the program still needs to confirm this. This takes 452 milliseconds.
3. The sending time is constant, about 158 ms, if the XBee had more than 2.5 seconds to join. However in case the XBee must send immediately after it is joined, the sending time is not constant and takes on average 611 milliseconds.
4. The sending time increases if there are more obstructions between the antennas.

With these characteristics we came to results discussed in the next sections. For the validation of these conclusions please see appendix A. Table 3.3 sums up the results of the distance-relation test.

Distance	Average sending time (ms)
Air	158
1 Floor	268
2 Floors	357
3 Floors	484
4 Floors	558
5 Floors	unreachable

TABLE 3.3: Distance consequence on send times

To save power the Waspmove can store the values for a user determined time. Taking samples and save them to EEPROM in case of hibernate mode takes only 6 - 7% of the time to measure and send. Table 3.4 confirms this.

Nr of samples per sensor	Average ON time (ms)
10	210
3	194

TABLE 3.4: Time needed to sample and store 4 sensors

3.4.3.2 Battery life with standard program optimizations

The application scenario for this battery test is as follow: the Waspmove will be turned on as short as possible and 4 sensors, namely temperature, humidity, pressure and battery level will be sampled. The node will take 10 samples for each sensor and calculate the average. Those values are put into one ZigBee packet and sent to the gateway. By adapting the sleep time between the event we came to the rather disappointing results shown in figure 3.9.

The graph in figure 3.10 breaks down the total energy consumption to five categories. It shows the monthly energy consumption as a function of the time between the events. For small intervals the active energy usage is huge. Only starting at 20 minutes sleep time the self-discharge becomes dominant and from 3 hours on the sleep mode current also becomes dominant.

Since the Waspmotes use this much energy when applied this way we will call this the High Performance mode from now on. The next section calculates an alternative approach, referred to as Power Saver mode. This nomenclature is continued in the program:

```
typedef enum {HIGHPERFORMANCE, POWERSAVER} PowerPlan;
```

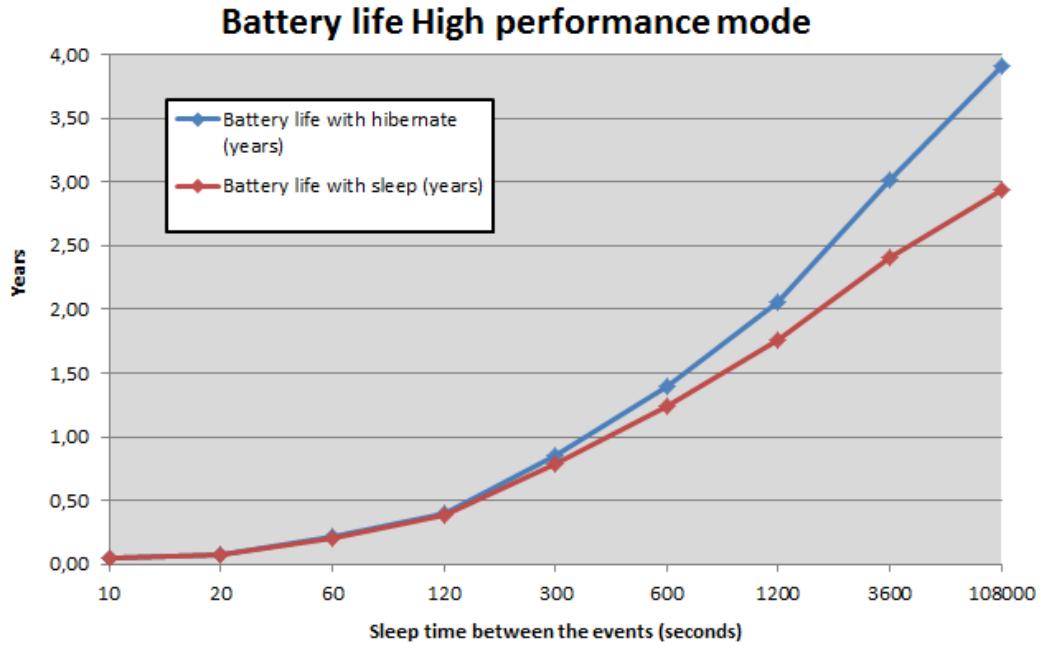


FIGURE 3.9: Battery life in High performance mode

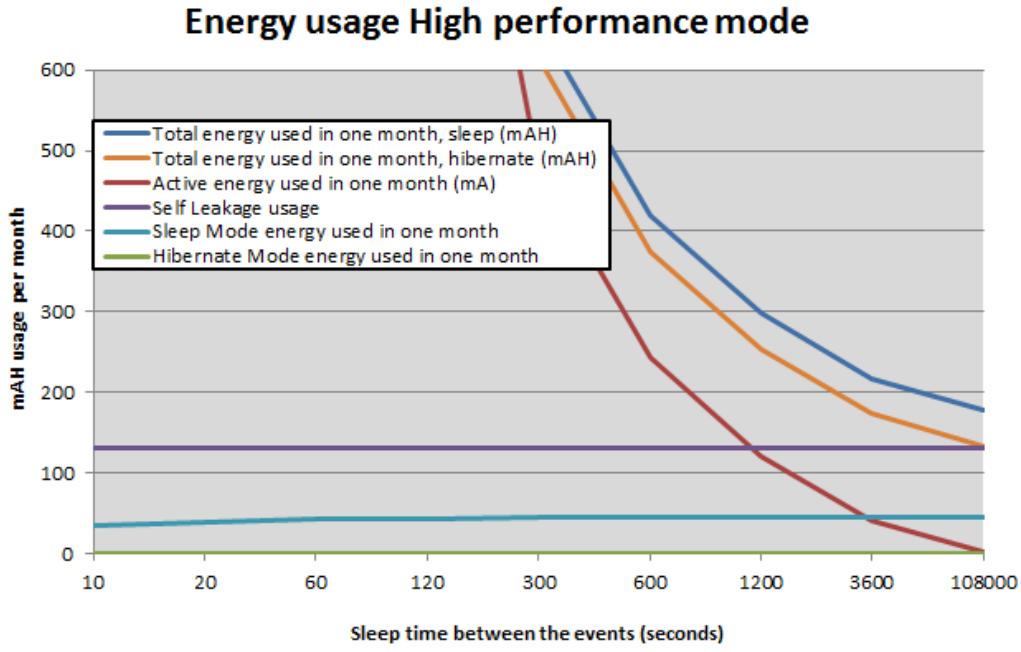


FIGURE 3.10: Energy usage in High performance mode

Table 3.5 summarizes the battery duration in years of both performance and power saver mode.

Sleep duration	Deep Sleep		Hibernate	
	High Performance	Power Saver	High Performance	Power Saver
10s	0,05	0,92	0,05	1,00
1min	0,21	2,15	0,21	2,63
3min	0,79	2,75	0,85	3,59
10min	1,25	2,85	1,39	3,76
20min	1,75	2,90	2,06	3,85
1h	2,41	2,94	3,02	3,91
3h	2,94	2,96	3,90	3,94

TABLE 3.5: Battery life in years for High Performance and Power Saver

3.4.3.3 Battery life with extra optimizations

As shown earlier in this section, sending values requires that the Waspmote is on for at least 3 seconds. In addition the XBee uses about five times the energy of the Waspmote. For end devices it is obviously recommended to turn on the XBee as little as possible, within a user defined limit.

Figure 3.11 shows the same results as the application scenario discussed in section 3.4.3.2 and adds the results for a mode further referred to as Power Saver.

The implementation of this mode will depend on the nodes sleep settings. For *Deep Sleep* the values can simply be stored on the heap, but for *Hibernate* the values must be written to EEPROM.

Because of the size limit of a ZigBee packet we can store maximum 30 values and send them in one packet. However, if the sensor measuring interval is small the user can opt to store more values and send two or more packets after each other. The values for Power Saver in table 3.5 are of an example scenario that takes 60 measurements and then sends them in two packets to the gateway. It are also those results which are put in function of time in figure 3.11.

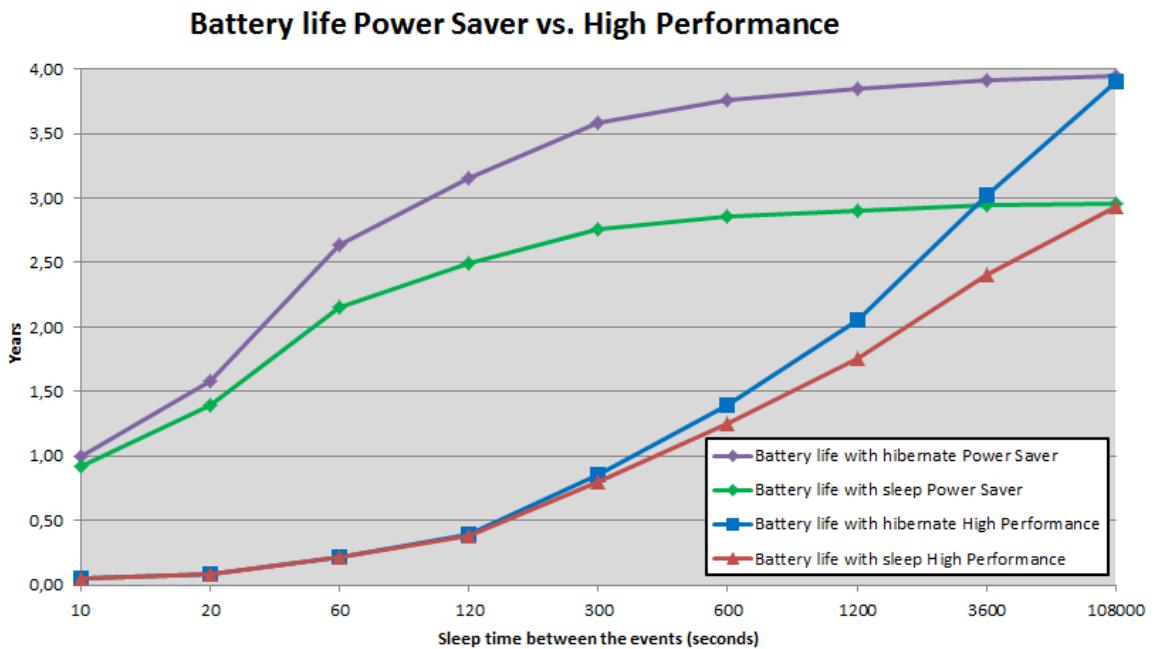


FIGURE 3.11: Battery life High Performance vs. Power Saver

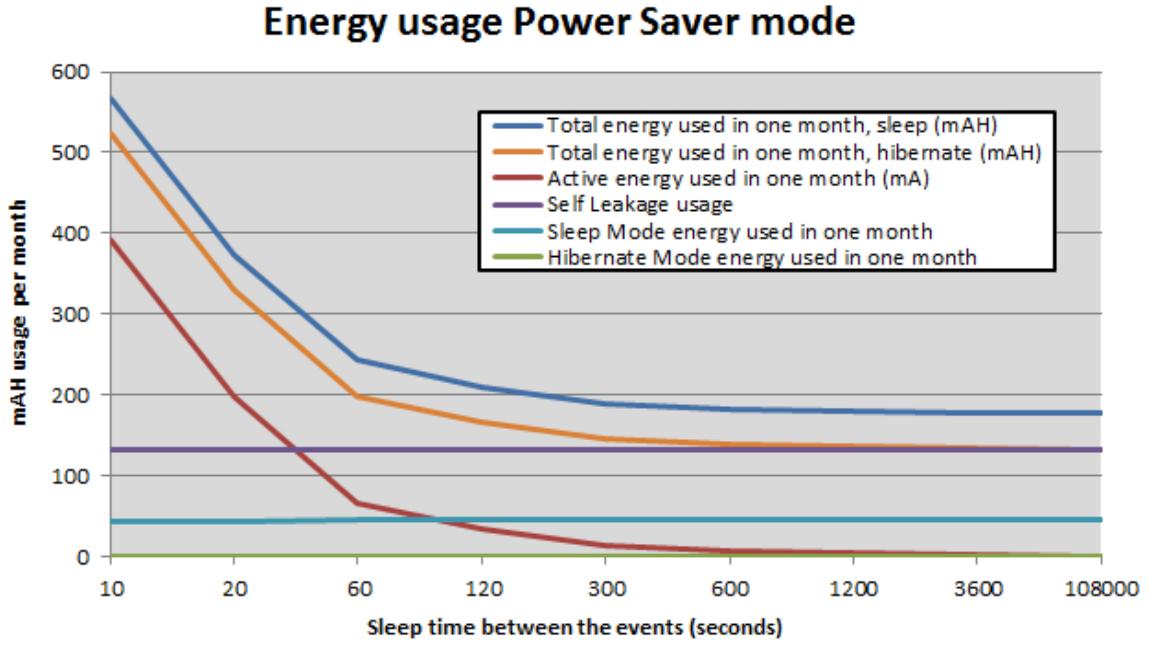


FIGURE 3.12: Energy usage in Power Saver mode

As visible on figure 3.11 *Hibernate* has more influence in Power Saver mode, already extending battery life significantly at a 20 seconds interval comparing to a 10 minutes interval in High Performance mode. Also the energy breakdown graph in figure 3.12 shows that the interval times must be increased much less before the dominant factor is self-discharge and sleep mode energy consumption, compared to figure 3.10.

By reducing the sensor measurement accuracy battery life can be extended with modest 3 - 4%, best case scenario. Please see appendix ?? for details.

In case the measuring intervals are small it is recommended to use *Deep Sleep* instead of *Hibernate*, since in hibernate the values are written to EEPROM. Equation 4.1 shows this can be very destructive for the Waspmote. Depending on how much freedom the user is given, the program can make the decision to switch to *Deep Sleep* on itself, or the installation's administrator can control this.

Chapter 4

A WSN with Waspmotes: Implementation aspects

4.1 Introduction

To start programming the Waspmotes the first step is to install the Waspmote-IDE. This IDE uses the same compiler (AVR) and core libraries as the Arduino IDE. According to Libelium their IDE has been properly tested and proven to assure optimum operation. Unfortunately we cannot agree with this. The Waspmote-IDE also includes Libelium API libraries to help you creating your Wasp mote programs. Libelium also offers a lot of program examples on their website. Sadly most of them don't work like Libelium claims they do.

4.2 Libelium IDE and API

4.2.1 Waspmote-IDE

The Libelium IDE offers some advantages compared to using other IDE's. For example by using Eclipse it is possible to update programs that are too big (*i* 120KB) for the Wasp mote, over-writing the bootloader. Then they must be sent back to Libelium to restore them. The Wasp mote IDE does not allow this accident, so Libelium does not support using other IDE's in an official way so that there is no valid warranty if you've erased the bootloader. Their IDE is far from perfect however, some issues we've experienced are:

- Opening a second or more instance of the IDE sometimes re-opens the previously active files, making it confusing to detect which one you were working in so you end up with two unsaved versions of the same code.
- Once you start compiling (which takes a lot of time) there's no way to stop it, the stop button does not work.
- Uploading immediately after compiling will first re-compile it anyway.
- There is no complete C/C++ support. For example using simple enums is not possible. A workaround is to place the code in additional .h or .cpp files.
- Auto-completion for the Libelium API functionality would be a great addition.

Also the Waspmove's (V1.1) hardware slows down the programming process:

- Uploading the code takes a lot of time: 1.5 - 2 minutes.
- The uploading process fails if:
 - The XBee is present
 - The hibernate jumper is not present when the mote is in hibernate
 - The little power switch has been turned off

Often you will want to turn off the power switch temporarily to analyse the content of the serial monitor. Especially in pair programming there is often one requirement you forget and the Waspmove does not check for this on beforehand. It will first compile and do as if it is uploading your code, disappointing you at the end of the process.

When debugging bigger program these actions come even more annoying. Suppose you are testing a program which measures sensors, sends the values and hibernates. Then you must:

1. Remove the sensor board
2. Place the hibernate jumper
3. Remove the XBee
4. Upload
5. Place the XBee
6. Remove the hibernate jumper
7. Re-mount the sensor board

And this is not the end of the list. Removing the hibernate jumper causes the Waspmove to crash one out of two times. Resetting the mote has no effect in this case, just keep inserting and removing the little jumper until it agrees with what you want.

4.2.2 Waspmove-API

To facilitate the programming Libelium offers a quite big API and after some exploring you are quickly started with it. The structure of the API is very simple, there is little to no inheritance. In this section, the most important classes are indicated with a red box and are recommended to explore before starting future programming.

4.2.2.1 Original structure

Each module or concept has its own class, for example all RTC functions are in 'WaspRTC.h' and 'WaspRTC.cpp'. To be able to use those functions in the IDE an object from the class 'WaspRTC' must be created. This object is created by default by the library and it is public to all libraries. All types to be run on the API can be found in the 'WaspClasses.h' file and each file also includes this file so it is aware of all available types. Please see appendix ?? for a complete overview.

For our application WaspXBeeZB is one of the most important classes. It inherits from WaspXBee Core and this way it is also related to the WaspXBee class. Figure 4.1 indicates this relation. Figure 4.2 displays the relationship with the more basic Libelium/Arduino libraries and the Waspmove hardware. For example `typedef unsigned char uint8_t` can be found in 'stdint.h'.

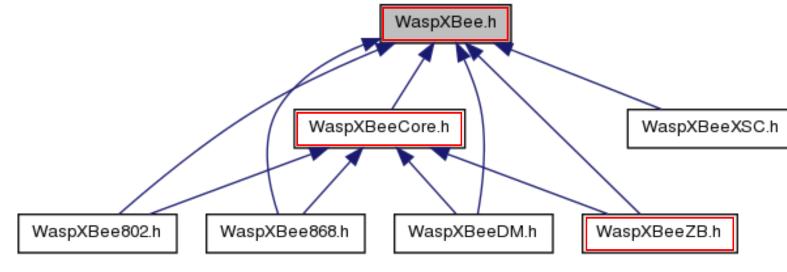


FIGURE 4.1: Reduced dependency graph for WaspXBeeZB

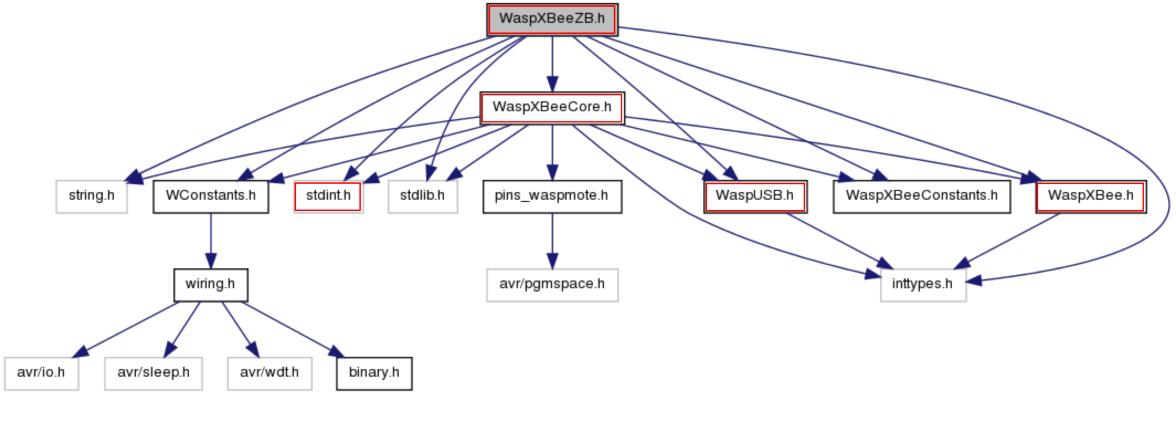


FIGURE 4.2: Reduced dependency of core libraries

During the development of our program the Waspmotet showed several strange effects that could only be explained by bad stack management or heap and stack conflicts. Because of this lack of free memory (SRAM, 8KB) we discovered that the V1.1 API wastes a lot of memory by always including all libraries despite not using them. As a fix Libelium recommends to remove all classes you do not use, and there fields that are used in other classes, 'just' going through the compiler errors one by one. After this our program had enough free memory and showed normal behaviour.

4.2.2.2 Added functionality

To facilitate programming extra functionality has been added to the Libelium API. They are inside files containing the original name with the 'Utils' addition, for example 'WaspRTCUtils.h' and can be found in 'BjornClasses.h'. To get an overview of the classes

for the practical part: memory (SRAM) problems: using = freeMemory(); Joe is true. V11 API wastes a lot of memory including always all libraries despite of you are not using them.

Try to remove every class you do not use, and also there are some certain cases that you have to remove the usage of some constants and functions.

just go one by one to the compiler msgs. We can help you if you are not sure removing something.

Regards I'm experiencing similar problems with my Waspmotets due to other issues, a recommendation is to remove from the API library, all that classes that you don't use such as SD.h and .cpp, or GPS, etc. bear in mind that you have to remove any reference from the file WaspClasses and where the compilator throws you an exception.

XCTU problems (cannot read / reset / writes always) cannot change device role via Libelium

4.2.3 Device start-up

Whether we come out of a hardware reset or a hibernate reset, the first thing the node will do is try to establish a connection with the network. Depending which reset we come from different routines will be executed. In normal operation mode this process only takes about 2 seconds (Please see appendix ?? for more measurement results). However, when a node is not able to join a network it will go into panic mode. This means that if the operating sleeping time is small, this time will be ignored and the node will wake up less frequently until it is able to rejoin the network, that way saving battery power. Supposing the node is able to rejoin the network, it will send the number of panics it experienced to the coordinator so a network administrator can investigate of the severity of the problem.

4.2.3.1 Full initialization

When the program is executed for the first time or when a hardware reset is detected the XBee will execute a full initialization process and the RTC will be set to zero. This means the default PAN ID and possibly other user settings like node ID will be written to the XBee. After this write the XBee must be re-setted (turn the power off and back on) and only than the joining attempts can start. This means full initialization takes about 9 seconds on average.

4.2.3.2 Reduced initialization

By not resetting the PAN ID but fetching it from the XBee's memory the joining process only takes about two seconds. Unfortunately a disadvantage of this shortened setup is that the XBee is no longer able to detect if the coordinator or his parent is actually available. The program will only notices this for the first time when it is trying to send a message. If this function results in a send error the program will do a full setup routine and resend the message. If the node then fails again to send the message we can conclude that the coordinator is really off-line or that there are no joinable nodes within range.

"Now, let's try to make a simple code which hibernates well. We recommend you to put a delay of a few seconds or a led blink before PWR.hibernate() sentence to allow removing hibernate jumper correctly."

4.3 Power savings

4.3.1 In the algorithm

4.3.2 Sleep vs. Hibernate

Compare battery life

4.3.3 Variable sleep times

In default configuration the Waspnote will send only its battery level to the default gateway and go into hibernate mode for 15 minutes. After this the cycle repeats. However it is possible that the node has more sensors implemented and the user wishes to obtain these values at different frequencies. In that case the node will have various different sleep times.

The next subsections will discuss different techniques to determine those sleeping intervals. In hibernation mode the node is completely disconnected from the main battery and the program stops. This makes that all variables lose their values and must be stored in EEPROM memory

if they must be known during the next cycle. Each of the next techniques present with benefits and drawbacks and since we are working with embedded systems with limited possibilities, one should also consider to limit the users options to facilitate the calculations.

4.3.3.1 Calculate only the next time to sleep

Each algorithm will have to store the individual sleep times per sensor. To support this algorithm also a copy of the original time will be saved and each time the node wakes up it will look for the smallest next time to sleep. This number will be subtracted from the other sleep times in the array. When a value becomes zero it will be restored with its original value and the cycle continues. The following example demonstrates the process:

This process is fast and simple. However, the main advantage is that the node has to write to

Sensor[4]	Sensor[3]	Sensor[2]	Sensor[1]	Sensor[0]
100	50	35	10	20

TABLE 4.1: Individual Sensor Sleep Times in seconds

Cycle	Sensor[4]	Sensor[3]	Sensor[2]	Sensor[1]	Sensor[0]	Sleep time
0	100	50	35	10	20	10
1	90	40	25	10	10	10
2	80	30	15	10	20	10
3	70	20	5	10	10	5
4	65	15	25	5	5	5
5	60	10	20	10	20	10
6	50	50	10	10	10	10

TABLE 4.2: Example of sleep algorithm 1

EEPROM each time it wakes up. According to Atmel the EEPROM of the ATmega1281 has an endurance of at least 100,000 write / erase cycles. The following equation indicates the problem for an interval of 10 seconds:

$$\frac{100000 \text{ writes} \cdot 10\text{s}}{60\text{s} \cdot 60\text{min} \cdot 24\text{h}} = 11,57 \text{ days} \quad (4.1)$$

But the processors has 4Kbytes EEPROM on board so we don't have to write to the same place every time. Since EEPROM is written on a 'per cell' basis this can extend the lifetime. Our sensor mask can contain up to 16 values of 2 bytes. This leads to the next result:

$$\frac{100000 \text{ writes} \cdot 10\text{s} \cdot 4\text{KB}}{60\text{s} \cdot 60\text{min} \cdot 24\text{h} \cdot 365\text{days} \cdot 32\text{B}} = 3,96 \text{ years} \quad (4.2)$$

We still must store where the data is stored but this won't cause big problems since we only have to rewrite this cell 125 times:

$$\frac{4\text{KB}}{32\text{B}} = 125 \text{ writes} \quad (4.3)$$

4.3.3.2 Calculate all next times to sleep

Another possibility is calculate as much as possible or if maybe even all sleep times.

4.3.3.3 Limit user control

4.3.4 Overdrive function

The overdrive effect amplifies the audio signal in a non-linear way. There are different possible schemes, we used the three layer non-linear soft saturation scheme below. The threshold value is set at 1/3, given that the input audio values are between -1 and 1.

$$f(x) = \begin{cases} 2x & \text{if } 0 \leq x < \frac{1}{3} \\ \frac{3-(2-3x)^2}{3} & \text{if } \frac{1}{3} \leq x < \frac{2}{3} \\ 1 & \text{if } \frac{2}{3} \leq x \leq 1 \end{cases}$$

- In the lower third the output is linear - multiplied by 2. for $\frac{2}{3} \leq x \leq 1$
- In the middle third there is a non-linear (quadratic) output response
- Above $\frac{2}{3}$ the output is set to 1.

Image 4.3 below shows the effect of the overdrive on an audio sample shown in the time domain. The blue signal is the original, the red signal is overdriven signal. (Time on the x-axis, amplitude on the y-axis.) Bringing this non-linearity in the signal is called clipping. Clipping is a process

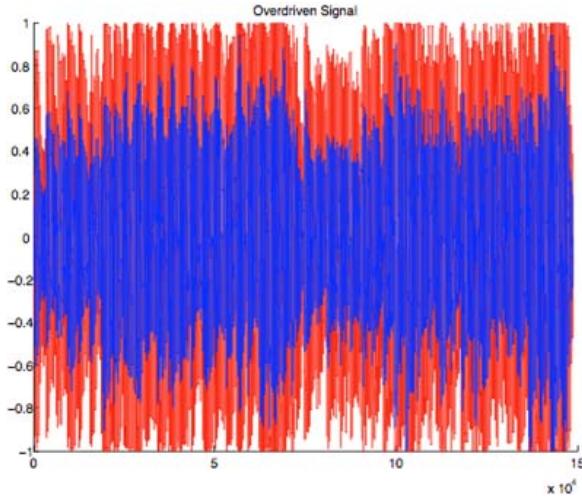


FIGURE 4.3: Effect of overdrive in time domain

that produces frequencies not originally present in the audio signal. These frequencies can either be "harmonic", meaning they are whole number multiples of the signal's original frequencies, or "inharmonic", meaning dissonant odd-order overtones.

Setting the threshold value also determines which type of clipping is applied. Increasing the threshold value will result in a softer clipping. Whereas decreasing the threshold value will result in a harder clipping. When a threshold value of 1 is defined, there is no overdrive applied. The output remains the original signal.

The hardness or softness of the clipping matters. Hard clipping results when the output wave equals the input up/down to a certain level, then stays at the clipping level until the input drops below the clipping level again, giving perfectly flat tops and bottoms to the clipped output. Soft clipping has no abrupt clipping level, but gently rounds the top/bottom of the output wave so the waveform is "softly" rounded on top/bottom, not flat-topped.

Soft clipping emphasizes the lower-order harmonics, the third and fifth, etc. Hard clipping has a mix skewed to the higher order seventh and up harmonics, which are harsher sounding.

In the figure 4.4, a normal sinewave with amplitude 1 is shown after being processed by the overdrive effect with a threshold equal to 1. With a threshold of 0.5 there is a small soft clipping visible, see figure 4.5. Both images have time on the x-axis and amplitude on the y-axis.

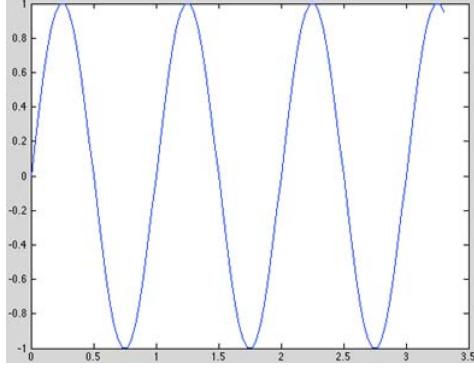


FIGURE 4.4: Overdrive threshold 1

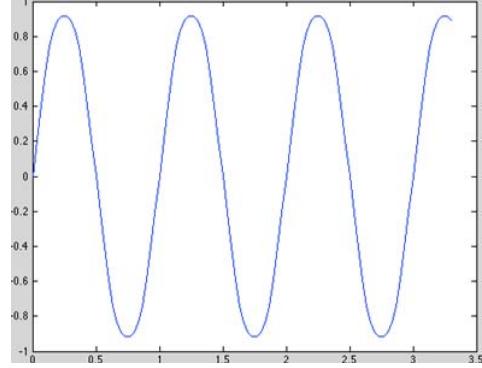


FIGURE 4.5: Overdrive threshold 0.5

4.3.4.1 Influence of the overdrive effect on the frequency spectrum

To test the effects of the overdrive effect on signals, a few examples are fed into a Fast Fourier Transform (FFT) application. We test the FFT with the following signal:

$$y(t) = 0.7\sin(2\pi \cdot 50t) + \sin(2\pi \cdot 120t) \quad (4.4)$$

This signal combines a sinewave of frequency 50 Hz and amplitude 0.7 with a sinewave of frequency 120 Hz and amplitude 1. When the FFT is finished and plotted, figure 4.6 below is the result. The graph shown in figure 4.7 the FFT of an overdrive with threshold equal to 0.1 on

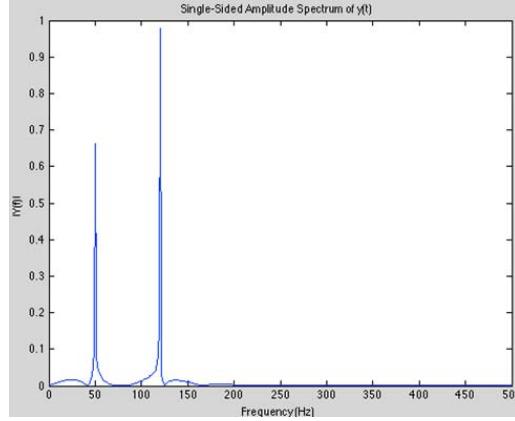


FIGURE 4.6: FFT of a basic signal

a sinewave with frequency 50hz. We can clearly see that there are only odd harmonics in the frequency spectrum, this is because the overdrive effect was set to symmetrical clipping.

4.3.4.2 Symmetrical clipping

For a given input waveform, say a sine wave, the tops and bottoms of the waveform are clipped equally, symmetrically. For a simple sine wave, symmetrical clipping generates only odd-order harmonics, giving a reedy, or raspy sound to the resultant waveform.

4.3.4.3 Asymmetrical clipping

The top(or bottom) of the waveform is clipped more than the bottom (top) half. This causes the generation of both even and odd harmonics, in contrast to symmetrical clipping's odd-order only. The even harmonics are smoother and more musical sounding, not as harsh as the odd ones. The hardness of the clipping and the degree of asymmetry affect the sound. The more asymmetrical, the more pronounced the even-order harmonics, the harsher the clipping, the more the harmonics are slewed toward higher orders.

The graph in figure 4.8 shows the Fourier transform of a the same signal, this time overdriven using asymmetrical clipping. Here the part of the sinewave above the x-axis was clipped using the same function as in the symmetrical clipping. The part of the sinewave below the x-axis is still the original. As we can see in the graph, using the asymmetrical clipping generates not only the odd but also the even harmonics. These even harmonics give an other type of sound.

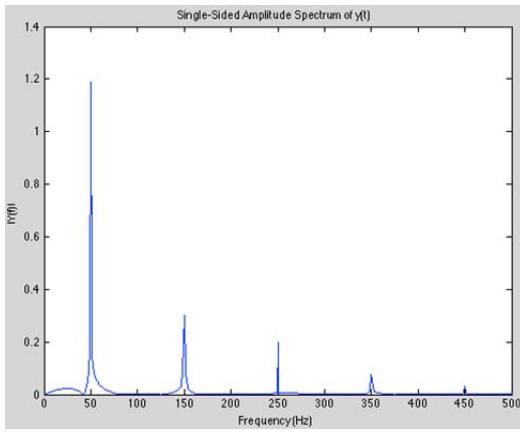


FIGURE 4.7: Symmetric clipping

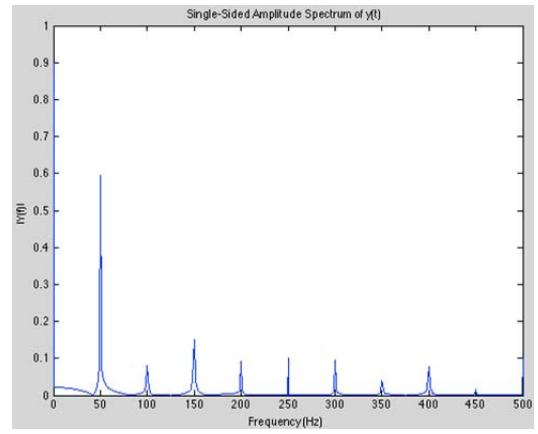


FIGURE 4.8: Asymmetric clipping

4.4 Equaliser

4.4.1 Design of the equaliser

The equaliser consists of three filters:

- Lowpass filter with cutoff frequency equal to 800Hz.
- Bandpass filter with cutoff frequencies equal to 800Hz and 6000Hz.
- Highpass filter with cutoff frequency equal to 6000Hz.

All the filters are constructed as FIR (Finite impulse response) with order 100. The filter components are created in MATLAB using the `fir1` command.

The architecture of these Finite impulse response filters is displayed in figure 4.9.

4.4.2 Amplitude response of the FIR filters

4.4.3 Testing of the equalizer in MATLAB

The equalizer has been tested by processing blocks of 1024 samples filled with white gaussian noise, the result has been fed into an FFT. Blocks are processed with the equalizer settings

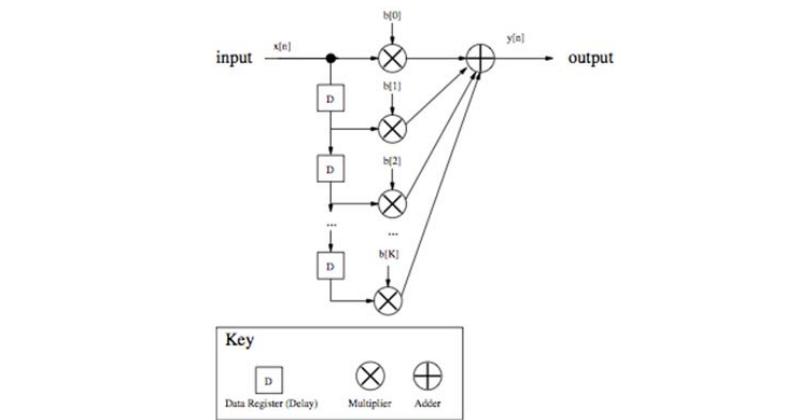


FIGURE 4.9: Architecture of the equaliser

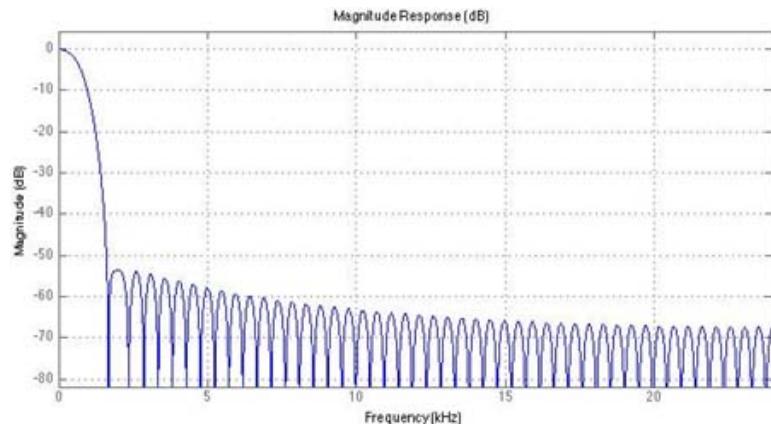


FIGURE 4.10: Magnitude response lowpass filter

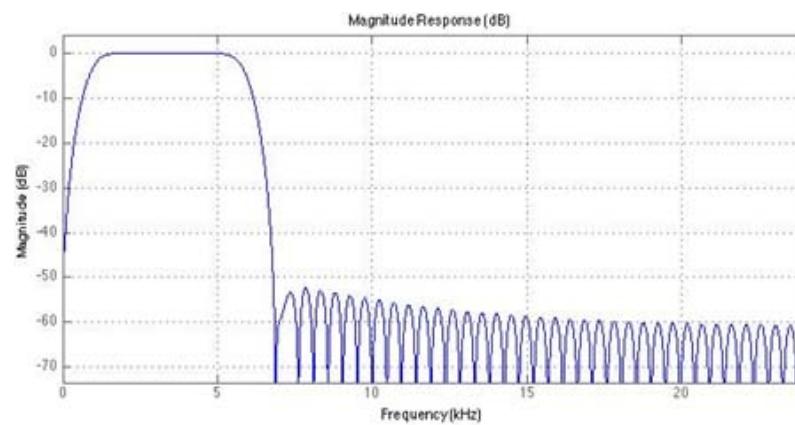


FIGURE 4.11: Magnitude response bandpass filter

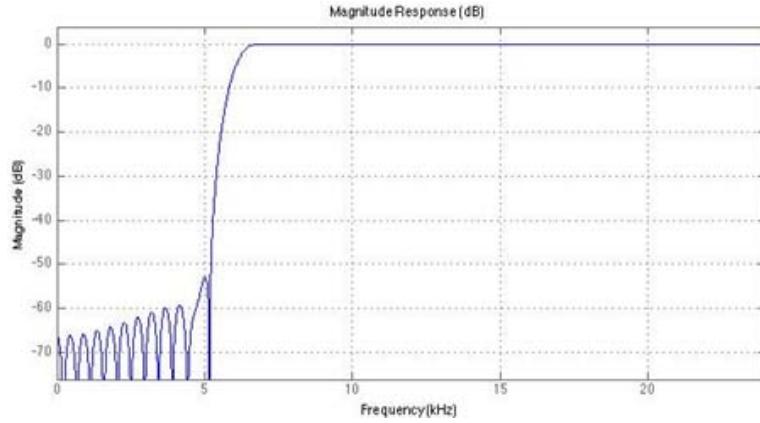
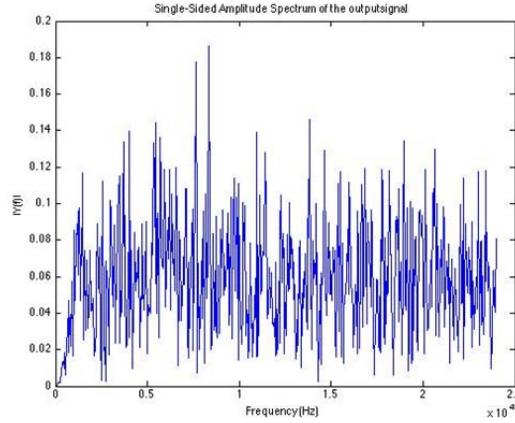


FIGURE 4.12: Magnitude response highpass filter

displayed in the caption of figures 4.13, 4.14 and 4.15.

The equaliser settings can be changed in the interface, using the sliders for all the filters where a

FIGURE 4.13: $\text{Gain}_{LP} = 0$, $\text{Gain}_{BP} = 1$, $\text{Gain}_{HP} = 1$

value between 0 and 10 can be set. The interface will respond by sending separate gain integers to the dsp board. Here every filter is independently multiplied by it's gain and devided by 10. Afterwards all the filter coefficients are added together to deliver the 101 final coefficients for the selected equaliser setting. These values will than be used to process the sample blocks entering the equaliser.

This is not the traditional way of FIR filtering, normally different coefficients are created for every specific equalizer setting. This way of working saves lots of processing power.

Because the FIR filters work with samples 'from the past', every sample block will first be prepadded with the 100 last samples of the previous block. These are stored in the so called prepad array. After the incomming block is prepadded, it 100 last samples will be stored in the prepad array. In this way always 1024 samples are loaded into the filer, this filter uses 1124 samples to create a 1024 sample outputblock.

4.4.4 Problems and solutions

- **Problem:** The overdrive effect can only generate odd harmonics.

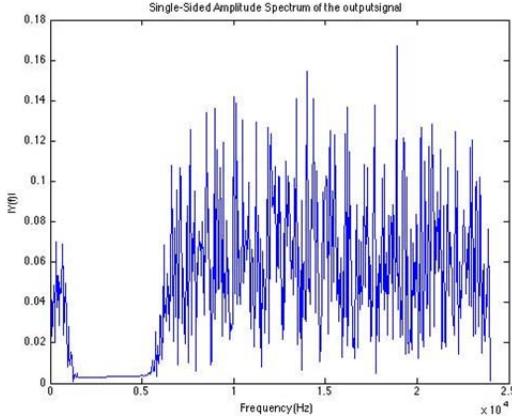


FIGURE 4.14: $\text{Gain}_{LP} = 1$,
 $\text{Gain}_{BP} = 0$, $\text{Gain}_{HP} = 1$

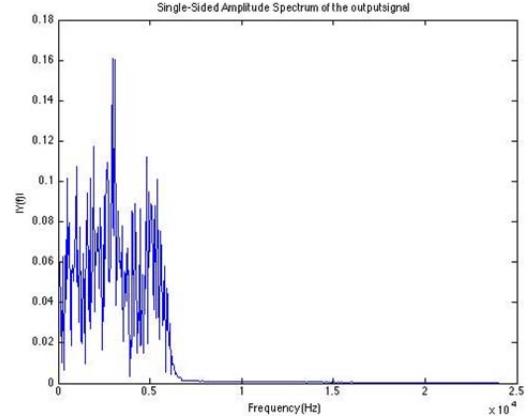


FIGURE 4.15: $\text{Gain}_{LP} = 1$,
 $\text{Gain}_{BP} = 1$, $\text{Gain}_{HP} = 0$

- **Solution:** An extra option is added, which lets the user choose between 'symclip' (odd harmonics) or 'asymclip' (even harmonics).
- **Problem:** The interface was not able to generate specific filter coefficients for all the different equalizer settings.
 - **Solution:** Solved by designing a filtering method which can use 'static' coefficients. Here the interface only has to send the different gain integers to the dsp board.
- **Problem:** The FIR filter goes back 100 samples in the past to process the current sample.
 - **Solution:** An extra 'prepad' array is added, initiated by zeros when the equalizer starts, and always saves the last 100 samples of the 1024 sample block. Allowing the filter to use these samples.
- **Problem:** The filter coefficients for the highpass filter are too large for the 32-bit double values, the dsp board uses.
 - **Solution:** We first changed the compiler settings to use 64-bit memory for the values, but this caused the rest of the program to stop working. It was finally solved by creating new filter coefficients for the high pass filter who were small enough to fit in 32 bits.
- **Problem:** The equalizer added an extra noise to the signal
 - **Solution:** The problem here was that the prepad array was not filled correctly, so it prepadded every sample block with 100 zeros as it was initiated. The last 100 samples were always saved in a wrong place in the memory. The use of pointers for this storing and fetching caused the program not to crash during runtime. Correcting the save location solved this.
- **Problem:** The bandpass filter acts as a highpass from its first cutoff frequency
 - **Solution:** at implement

Chapter 5

Optimizing Wireless Sensor Networks

Written by Bjorn Deraeve

5.1 Introduction

5.1.1 What is digital signal processing?

Signal processing falls within the scope of electrical engineering and applied mathematics that deals with the analysis of signals or operations on signals. In order to do this signals are presented in discrete time, discrete frequency or other discrete signal domains. The set of algorithmic solutions that are used to process digital signals are limited by the processing capabilities of the available hardware. Digital Signal Processing (DSP) algorithms have long been run on standard computers. Today additional technologies such as specialized processors called Digital Signal Processors (DSP) are used. DSP processors can perform typical DSP operations more efficiently thanks to single-cycle multiply-accumulate (MAC) units, shift registers and extra large accumulators. To increase speed and possible DSP realizations DSP processors are used in purpose-built hardware such as application-specific integrated circuits (ASICs) or field-programmable gate arrays (FPGAs).

5.1.2 How are DSP processors used?

DSPs receive real-world signals like voice, audio, video, ... that have been digitized and then mathematically manipulate them. DSP processors are a special type of microprocessors which are designed for performing mathematical functions like "add", "subtract", "multiply" and "divide" very quickly. Real-world analog signals are converted by Analog-to-Digital converters and are turned into the digital format of 0s and 1s. Now the DSP can start working on this digitized information. When all processing is done the new samples are mostly feed back in the real world via Digital-to-Analog converters. This whole process happens at very high speeds.

For our FM synthesizer application the first step is different. The DSP processing algorithms consist of algorithms that create the signal and signals that manipulate those signals so the AD converters on our FPGA are not used.

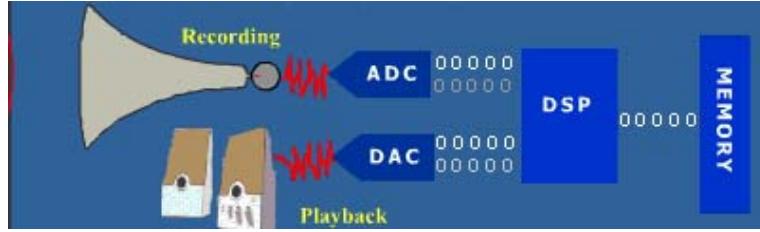


FIGURE 5.1: DSP Principle

5.2 Digitization

Real-world continuous signals have to be reduced to discrete signals (a numeric sequence) before a DSP can manipulate them. This sampling is done by an analog-to-digital converter (ADC) and the resulting samples are transported serially to the DSP.

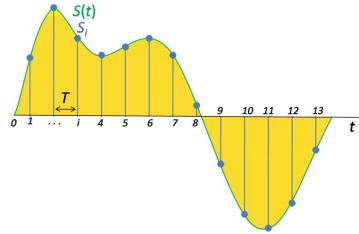


FIGURE 5.2: Signal sampling representation

The sampling frequency f_s is the number of samples taken in one second. According to the Nyquist criterion, the sampling frequency must at least be twice the frequency of the highest frequency component of interest. If the sampling frequency is greater than or equal to twice the bandwidth of a bandlimited signal then the signal can be (re)constructed without aliasing. This frequency is called the Nyquist rate.

5.3 Digital Signal Processors

As briefly mentioned in this chapter's introduction DSPs are processors with hardware, software, instruction sets, parallelism and data addressing that are optimized for high-speed numeric operations. A DSP contains the following key components:

- **Program Memory:** here the processing algorithms are stored
- **Data Memory:** stores the signals to be processed
- **Compute Engine:** performs the mathematical operations, coordinates access to program and data memory
- **Input/Output:** connections to the outside world

The need for high speed processors is because most DSP calculations are done on real-time signals. The input signal comes to the DSP as a train of individual samples from the AD converter. To do for example filtering in real-time, all calculations and operations on a sample must be completed before the next sample arrives. Usually also previous samples are needed so these must be stored somewhere and have to be quickly accessible.

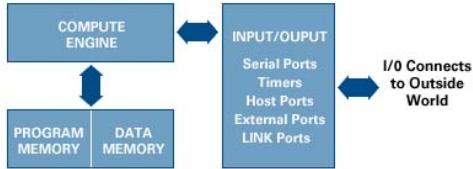


FIGURE 5.3: DSP structure

5.3.1 Numeric architecture

We know that DSPs must complete multiply-accumulate, additions and bit-shift operations in a single instruction cycle. Hardware optimized for such numeric operations is typical to DSP processors and distinguishes them from other general-purpose microprocessors. The numeric operations are done by the DSP's MAC units, the arithmetic-logic unit (ALU) and barrel shifters:

- **MAC:** performs sum-of-products operations (used by FIR and IIR filters and fast Fourier transforms). The multiply-accumulate operation computes the product of two numbers and adds it to an accumulator:

$$a \leftarrow a + (b \times c) \quad (5.1)$$

The multiplier is typically implemented in combinational logic followed by an adder and finally the accumulator register to store the result. If the output of the register is fed back to the adder each clock cycle the output of the multiplier is added to the register. Calculating multiplications with combinational logic requires a large amount of logic but is much faster than the method that uses shifts and additions.

- **ALU:** performs standard addition, subtraction and basic logical operations
- **Barrel shifter:** performs operations on bits and words. This digital circuitry can shift a word by a specified number of bits in one clock cycle and is implemented as a sequence of multiplexers where the output of one mux is the input of the next mux. The input depends on the shift distance. An example of an 8-bit barrel shifter is shown in Figure 5.4.

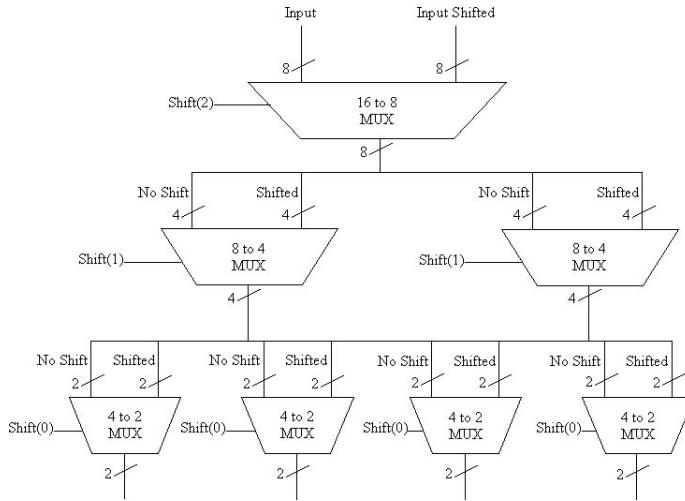


FIGURE 5.4: 8-bit barrel shifter

A barrel shifter is for example used for the implementation of floating-point arithmetic. To add two floats their significands must be aligned, which requires shifting the smaller number until it matches the exponent of the larger number.

As mentioned also parallelism plays a role in the efficiency of DSP processors. Figure 5.5 shows the parallelism of the ALU, MAC and shifter.

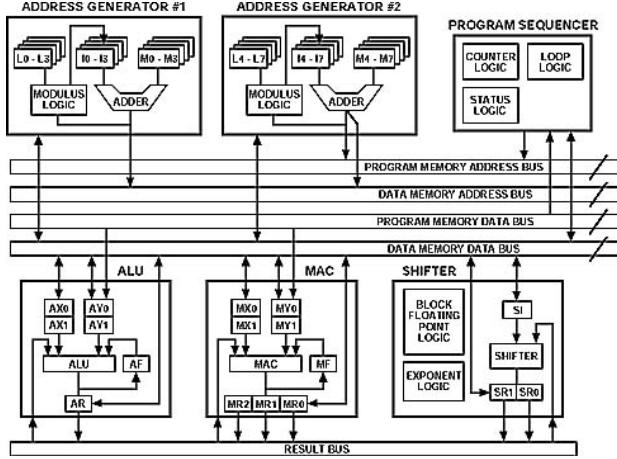


FIGURE 5.5: Useful DSP architecture

5.3.2 Memory architecture

Of course the memory of a DSP processor must be able to follow the high speed of the operations. Bus architecture must be optimized, there cannot be delays or bottlenecks.

- Most general-purpose microprocessors use the von Neumann architecture and throughput is limited because of having to choose between either fetching data or fetching an instruction in each cycle. In DSP processors program memory and data memory have their own space and busses and can both be fetched in one cycle, doubling throughput. This structure is known as the Harvard architecture. Naturally additional optimizations present on general-purpose processors, like instruction cache for example, are also present on DSPs.

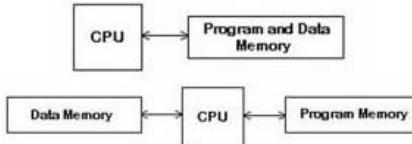


FIGURE 5.6: Memory architecture: Von Neumann (top) vs. Harvard (below)

- A lot of DSP algorithms need to get data from memory in repeating patterns. For example the fetching and manipulation of samples stored in an array. By reducing instructions needed for memory access instruction cycles can be saved. To reduce this overhead DSPs use specialized data address-generators (DAG) to automatically manage these types of repeated memory accesses. Because most DSP algorithms use two operands the DSP has two DAGs. One address generator creates the address for the data memory bus, the other one for the program memory bus. By this way the DSP is able to sustain execution of instructions in one single cycle.
- Often DSP algorithms require data in a range of addresses going from the end of the buffer back to the start of the buffer (wrap around). Therefore the DSP is provided with hardware circular buffers. The DAGs are specially designed for this task. General-purpose processors perform these functions in software and limit the ability to handle real-time signals.

5.3.3 Sequencer architecture

Since most algorithms run on a DSP processor are by nature repetitive, the compute engine's program counter must be able to loop through the code without overhead while getting from the end of the loop back to the start. In general-purpose CPUs the loop's end condition is fully maintained in software. This condition instruction requires the fetching of addresses from memory and evaluation and takes time (cycles). DSP processors perform these tests in hardware, storing the needed addresses. In one cycle the condition is evaluated and either the first instruction of the loop is executed or the first instruction outside the loop is executed.

5.3.4 Input/Output architecture

Because of the high data throughput requirement of the DSP the whole design is focused on making data accessible for the numeric, memory and sequencer sections. To transfer data into and out of the DSP quickly serial communications and DMA is used. The use of serial communications above parallel is naturally since this way data is sent one bit at a time, following the DSPs operating frequency and distributing work load. So the DSP communicates with ADCs, DACs or other devices through synchronous serial ports (SPORT). Other functions such as timers and boot logic (in other words: hardware interrupts) also ease DSP system design.

5.4 Fixed-point vs. floating-point DSP

Digital signal processing can either be done with fixed point or with floating point data storage. DSPs designed to represent and manipulate integers are fixed point realizations. Floating-point DSPs represent and manipulate rational numbers. The number is represented in a similar manner to scientific notation, with a mantissa and an exponent (e.g. $A \times 2^B$, where 'A' is the mantissa and 'B' is the exponent).

5.4.1 Advantages of floating-point digital signal processors

In a fixed point DSP numbers are represented with a fixed number of digits after the decimal point, for example: 123.45, 1234.56, 12345.67, etc. In a floating point implementation the decimal point can 'float'. In addition a floating point representation can represent 1.234567, 123456.7, 0.001234567, etc. Thanks to the exponential representation floating-point DSPs support a much larger dynamic range: very small numbers and very large numbers can be stored. Another advantage of this is that floating-point processing yields much higher precision than fixed point processing.

5.4.2 Considerations

In digital signal processing rounding and/or truncating numbers yields quantization error or noise, a difference between the real-world value and the quantized digital values. This also happens each time a DSP does a mathematical calculation. As such, floating-point processors are the ideal DSP when accuracy is critical, like in audio applications. This is noticeable in the program of our FM synthesizer, where constantly calculations are done with float data. When choosing between fixed-point or floating-point processor there are many other important factors to consider, e.g. processor cost, ease of development, performance. In summary, because of the greater number of general purpose applications that can use fixed point processors those are typically less expensive than floating-point versions due to the scale of manufacturing. Floating-point DSPs are optimized for computationally intensive applications.

Chapter 6

Implementation: C

Written by Bjorn Deraeve

6.1 Introduction

In chapter 5 we already looked at the DSP architecture and its advantages compared to traditional CPU's. Now we will have a closer look at programming concerns unique to DSP algorithms and some related hardware provisions.

More specifically real-time aspects will be discussed. Obviously in analog systems every task is performed in real-time. All signals and processing are continuous. However in digital systems signals are discrete and represented by a set of samples. The interpretation of real-time is depending on the sampling rate. In order for a digital system to be operating in real-time, all processing of a given amount of data must be completed before the next amount of data arrives. Thus real-time is limited by the amount of data to process and by the type of processing that must be done on that data. These two constraints correspond to the sample frequency and the complexity of the running algorithms.

6.2 Time management

Since there is only a limited amount of time to perform any given algorithm, time management is an essential part of DSP software design. The resulting strategies determine how the processor is notified about events, how data is handled and how communication is done.

6.2.1 Interrupts

Event handling by interrupts logically is the most appropriate way to meet to the timing constraints. Naturally the interrupt rate is based on the data sampling rate. Typically DSP processors have extra built-in hardware mechanisms to handle these interrupts because they are more efficacious than software. Interrupt service routines (ISR) on a DSP processor must comply to all of the following demands:

Fast context switching: can be accomplished by using a second set of data registers. Only one set is active at a time, containing all the relevant data for the active context. When an ISR is called the switch can take place immediately without first having to temporally save the active registers.

Nested interrupt handling: to handle nested interrupts, DSPs record their state corresponding to a certain context on a stack located in the DSPs program sequencer when an event occurs. This data structure allows that interrupts with higher priority can interrupt one with lower priority.

Continue to accept data while the ISR is running: for this two hardware features are provided:

- Interrupt latch: makes sure that no important events can be missed while servicing an interrupt.
- Automated I/O: external devices can put data into the DSP's memory without requiring intervention from the processor, so no data is missed while the DSP is busy (serial ports, DMA, autobuffering).

6.2.2 Data handling

Interrupts are typically generated due to data flow and may either occur with each sample, or after a block of samples has been collected. Each DSP's serial port has two I/O registers: a receive (Rx) and a transmit (Tx) register. When a (serial) word is ready to transmit or be received an interrupt is generated. In the ISR either the value of register Rx is read into a data register or a value is copied from a data register to the SPORT's Tx register. Obviously if I/O have the same sampling frequency one interrupt suffices.

Real-time systems typically introduce processing delay. A delay is introduced due to the A/D and D/A converter, the execution of the algorithm on the samples and by block processing. The idea of block processing is to process the signal not sample-by-sample but in frames of n samples. Block processing allows that multiple operations are performed in parallel: samples are collected, samples are processed and samples are output. With this data handling technique performance can also be increased if the samples are transformed into the frequency domain and for example a Fast Fourier Transform (FFT) is executed.

6.2.3 Real-time interface issues

To create a real-time environment not only fast interrupt response, efficient data handling and fast program execution is important, but also the flow of data in and out the processor is significant.

However parallel transport is for the less obvious, it can also be used. In that case they are typically at least as wide as the processor's word width. Parallel transfer is always faster than serial transfer while this high speed is often not necessary. Each cycle data can be sent or received. This requires very fast peripherals to keep up with it (for example SRAM chips). Because this very high speed is mostly unnecessary and because of the higher cost mostly serial transfer is used.

Unlike the RS-232 protocol the serial ports (SPORT) of DSPs use a synchronous interface. The clock, receive data, transmit data and important frame (cf. block of samples) synchronization signals are passed through the interface and guarantee optimal communication. For example for left- and right-channel audio signals the serial data stream can also be time-division multiplexed (TDM).

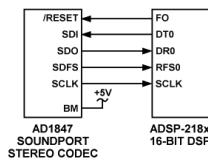


FIGURE 6.1: Serial communication interface

6.3 Analog Devices Sharc DSP Board

6.3.1 Introduction

Analog Devices has a wide portfolio of embededd processors, DSPs and microcontrollers for lots of general-purpose and application-specific goals. The SHARC processor we used dominates the floating-point DSP market. It's key features are summarized below. In appendix ?? a full functional block diagram for the ADSP-21369 is added for completeness.

Performance: 32-bit, 400MHz SIMD Core capable of 2.4 GFLOPS peak performance

32-bit external memory interface: supporting SDRAM, SRAM and FLASH

Digital Audio Interface (DAI): supporting access to user-definable peripherals like SPORTs, S/PDIF, precision clock generators, ... *In this project the DAI is used to output the generated signal to the audio jack.*

Digital Peripheral Interface (DPI): supporting access to user-definable peripherals like SPI ports, UARTs, I²C and timers. *In this project the DPI is used to communicate with the labview interface and will be further discussed in chapter 7.*

DMA: contains 34 zero-overhead DMA channels

6.3.2 Initializing communication between the DSP and the audio jack

6.3.2.1 Digital Applications Interface (DAI)

As mentionned above, the DAI is used to output our sound signal through the headphone jack. A simplified system architecture block diagram is showed in figure 6.2 and 6.3.

The DSP's DAI pins (DAI-P20-1) are connected to signal routing unit 1 (SRU) of the processor. The SRU is a flexible routing system which allows to route the DAI pin to different internal peripherals (SPORTs, S/PDIF, PCGs) in numerous combinations. Because the inputs and outputs of the peripherals are not directly connected to the processor but via the SRUs, an arbitrary amount of peripherals is possible which allows a wider variety of compactible applications without having to increase the pin count. So the SRU is simply a group of multiplexers under software control (memory-mapped registers). For our synthesizer application the SPORT peripheral is used.

The DAI is also connected to the AD1835 audio codec which makes in turn the connection to the audio jack, see figure 6.6.

6.3.2.2 Signal Routing Unit (SRC)

Like is said above, the SRUs are used to connect inputs and outputs. The signals routed by the SRUs are diveded into functional groups. Each group routes a set of signals with a specific purpose. For example, SRU1 Group B routes serial data signals. It are both the peripherals control registers and the SRU configuration that determine the direction of the signal flow in a pin buffer.

To configure the SRU we must thus write values that correspond to signal sources into the bit fields that further correspond to the signal inputs. However the registers are arranged into functional groups, this is stays a very difficult task. In order to ease this coding process one of the following two procedures can be used:

Expert DAI plug-in: This software plug-in is included in the VisualDSP++ IDE and greatly simplifies the task of connecting the signals in the SRUs.

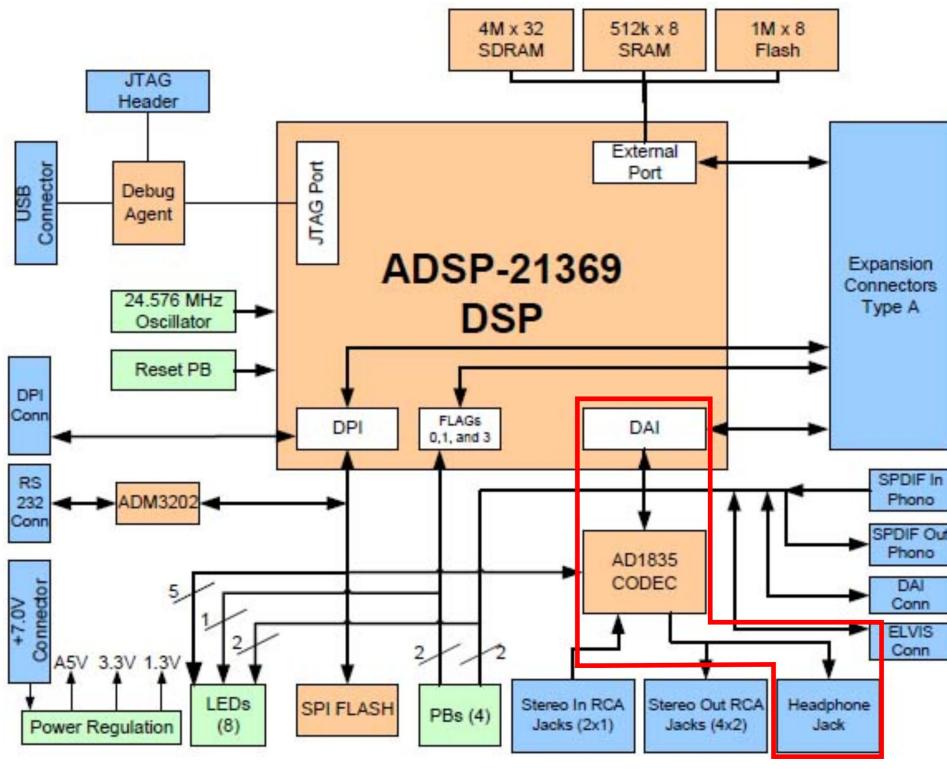


FIGURE 6.2: System Architecture Block Diagram

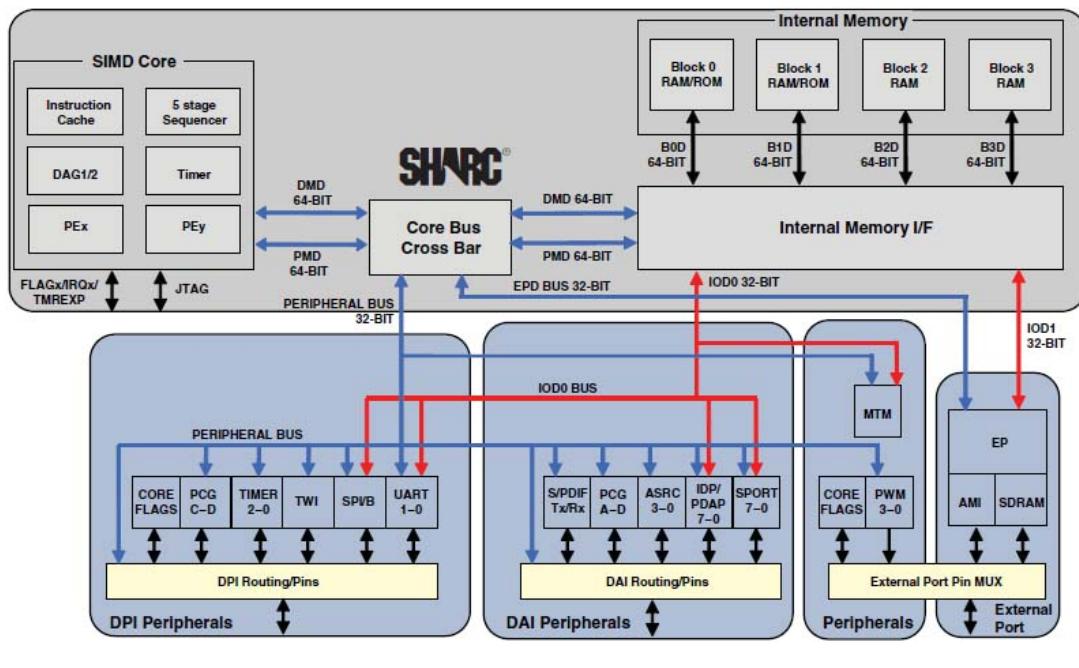


FIGURE 6.3: Peripherals Architecture Block Diagram

sru.h: This is a macro implementation that automates most of the work of signal assignments and functions. The macro can be used like is shown below to connect the ADC to SPORT0 using data input A. (See file *initSRU.c*)

```
#include <sru.h>
/* The following lines illustrate how the macro is used: */
/* Route SPORT 1 clock output to pin buffer 5 input */
SRU(DAI_PB05_0,SPORT0_DA_I);
```

To explain the input/output mnemonic, which always ends with `_I` if the signal is an input and with `_O` if the signal is output, we must explain one more thing in this section. Because of the context of the SRU, a pin is not just a pin anymore. Pins are now replaced by logical interfaces (pin buffers), shown in figure 6.4.

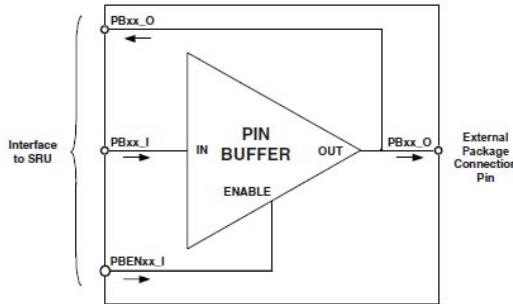


FIGURE 6.4: Pin Buffer

The interface's input must be thought of as the input to a buffer amplifier which can drive a load on the external pin. The pin buffer enable pin is an input signal that enables the buffer if its value is logic high and disables the buffer when its logic low. When the pin enable is set the output pin is logically equal to pin input. If not, the output of the buffer becomes high impedance and an external device can sent a logic value to the input of the ADSP-21369. The SRU receives this pin as an pin interface output and routes it as an input to the SHARC processor. If pin enable is asserted pin output is equal to pin input (the amplifier then acts as a current source, low impedance). See appendix ?? for more explanation.

6.3.2.3 AD1835A

The pins of the DAI are connected to an AD1835A audio codec. This is a high performance, single-chip codec with four stereo DACs and one stereo ADC, so 8 DACs and two ADCs. The engine has a programmable interpolator which allows the user to select different interpolation rates. By this way the desired sample rate can be selected. For the DACs the user can choose between 48kHz, 96kHz and one DAC even supports 192kHz, see the table below extracted from the AD1835A's datasheet. The location of this chip on the ADSP-21369 is shown in figure 6.5. Table 6.1 shows the connections between the processors DAI and the AD1835A audio chip. We used the 96kHz sample frequency.

Sample Rate	Interpolator Rate	DAC Control 1 Register
48 kHz	8x	000000xxxxxxxx00
96 kHz	4x	000000xxxxxxxx01
192 kHz	2x	000000xxxxxxxx10

TABLE 6.1: AD1835A DAC Sample rate settings

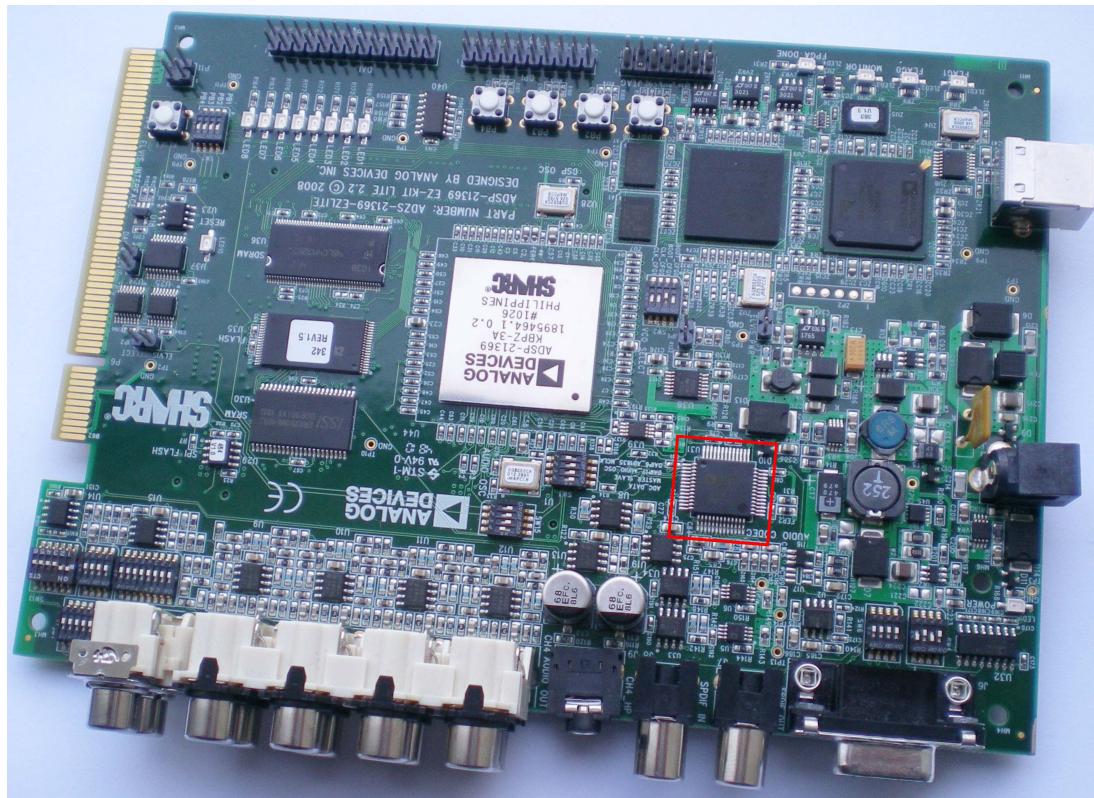


FIGURE 6.5: AD1835A on the ADSP-21369 DSP Board

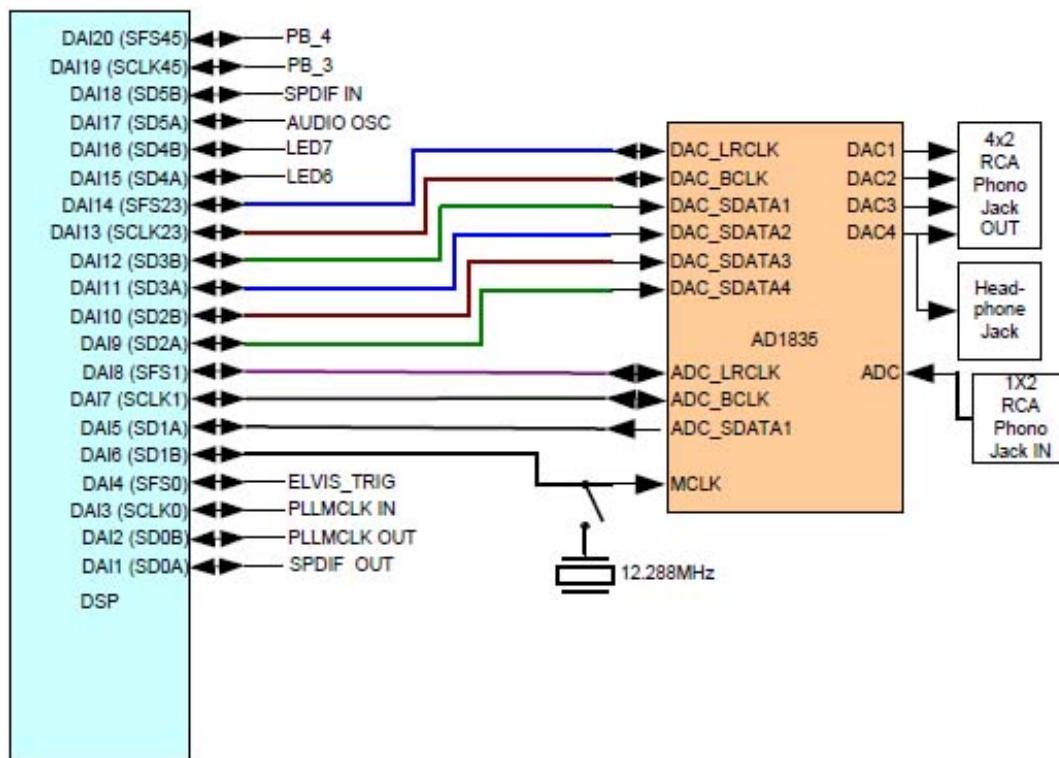


FIGURE 6.6: Peripherals Architecture Block Diagram

To make the AD1835A's four serial data output connections shown in figure 6.6 we use the macro defined in `<sru.h>`, which is shown below. For the complete initialization (clock and frame sync) we refer to the file `initSRU.c`.

```
SRU(SPORT2_DB_0,DAI_PB09_I);  <-- notice that DAC4 is connected to SPORT2B
SRU(SPORT2_DA_0,DAI_PB10_I);
SRU(SPORT1_DB_0,DAI_PB11_I);
SRU(SPORT1_DA_0,DAI_PB12_I);
```

6.3.2.4 SPORT

The previous sections explained how to set up the connections between the DAI and the DAC but we still have to initialize the communication protocol to be used. The ADSP-21369 has eight independent, synchronous serial ports (SPORT) that provide us with an I/O interface. Each SPORT has a set of control registers and data buffers which allows a variety of serial communication protocols.

In the synthesizer program we wil use SPORT2B to communicate from processor to DAC4 in I²C mode. The serial data will be transferred automatically from on-chip memory using DMA block transfers in a chained DMA process.

Chained DMA: In such a process the block to transfer comes from the (next) chain pointer register. When the current DMA transfer is complete the processor automatically begins the next DMA transfer indicated by the pointer register. To set up a chain of DMA operations the following steps are used:

1. Set up all TCBs in internal memory:

```
unsigned int Block_A[NUM_SAMPLES] ;
unsigned int Block_B[NUM_SAMPLES] ;
unsigned int Block_C[NUM_SAMPLES] ;

TCB_Block_A[0] = (int) TCB_Block_C + 3 - OFFSET + PCI ;
TCB_Block_A[3] = (unsigned int) Block_A - OFFSET ;

TCB_Block_B[0] = (int) TCB_Block_A + 3 - OFFSET + PCI ;
TCB_Block_B[3] = (unsigned int) Block_B - OFFSET ;

TCB_Block_C[0] = (int) TCB_Block_B + 3 - OFFSET + PCI ;
TCB_Block_C[3] = (unsigned int) Block_C - OFFSET ;
```

2. Set up the appropriate DMA control registers: set DMA enable bit SDEN_B, set chaining enable bit SCHEN_B. This information is written in the I²C mode control bits of the SPORT, SPCTLx.

```
/* Channel enable: SPEN_B
 * Word length: SLEN24
 * Operation mode enabled: OPMODE
 * Activate TXSPxy data buffers and transmit shift registers: SPTRAN
 */
*pSPCTL2 |= (SPTRAN | OPMODE | SLEN24 | SPEN_B | SCHEN_B | SDEN_B) ;
```

3. Write the adres of the first (next) TCB to the chain pointer register. For our program we use the SPORT Chain Pointer Registers CPSPx, so for SPORT2B this gives:

```
*pCPSP2B = (unsigned int) TCB_Block_C - OFFSET + 3 ;
```

For the complete initialization of the SPORT please have a look at file `initSPORT.c`. Remark that the data for the transmit buffers TXSPxA/B is loaded automatically by the DMA controller.

SPORT interrupt: The SPORT DMA mode provides a mechanism to transmit an entire block of serial data and creates an interrupt afterwards rather than after each transmitted word, significantly reducing overhead.

As is explained in section 6.1 all processing of the algorithms must be done before the next block of data is provided. To keep track of this real-time constraint the semaphores *isProcessing* and *blockReady* are introduced and update the real-time flag each time a SPORT DMA interrupt occurs. If the program becomes too slow, real-time is set false and a warning sound will be produced. If the processing is done before the interrupt occurs the SPORT chain pointer register is updated:

```
void TalkThroughISR()
{
    if(isProcessing) //still busy calculating
        realtime=0;

    (int_cntr++); //Increment the block pointer
    int_cntr %= 3;

    blockReady = 1;
}
```

6.4 The DSP Synthesizer program

All instructions for the Sharc processor are written in C or assembly. Once finished all source code is translated and linked by the compiler integrated in VisualDSP++. Then this executable machine code is uploaded to the device.

6.4.1 Program execution

To follow the next descriptions please take a look at figure 6.7. The *main.c* file contains the central starting point, namely the *main()* function. At run time, the execution will start at the first line of this function. In our case this is a series of *Init()* functions, which were briefly discussed in the previous section.

Once all initializations are done the real work can be started. At this point we enter the *while(!stop)* loop. As long as the user does not press the stop button the microprocessor continues running and must do something, so an infinite loop keeps it busy. This is typically for embedded controllers.

Next is another while loop: *while(blockReady)*. After the previous block has been processed and the SPORT DMA interrupt has occurred the next block can be processed. This next block will either produce the pursuit of the previous frame or will produce a warning tone, indicating that the calculations on the previous block were not finished before the block was needed by the DAC, and thus the signal could not be produced in realtime. We will first explain the not-realtime case since this will also be helpful in the realtime process.

6.4.1.1 Not-realtime

The not-realtime-function takes two parameters as argument, both pointers and creates a 8000Hz warning tone. The first pointer is the address of the array containing the samples to be processed. This array contains 1024 samples and is contained in a TCB block, used by the DMA chain to feed the DAC. The second is a simple int pointer which is used to indicate which particular samples to create, fulfilling the role as the time of the sine function in this case.

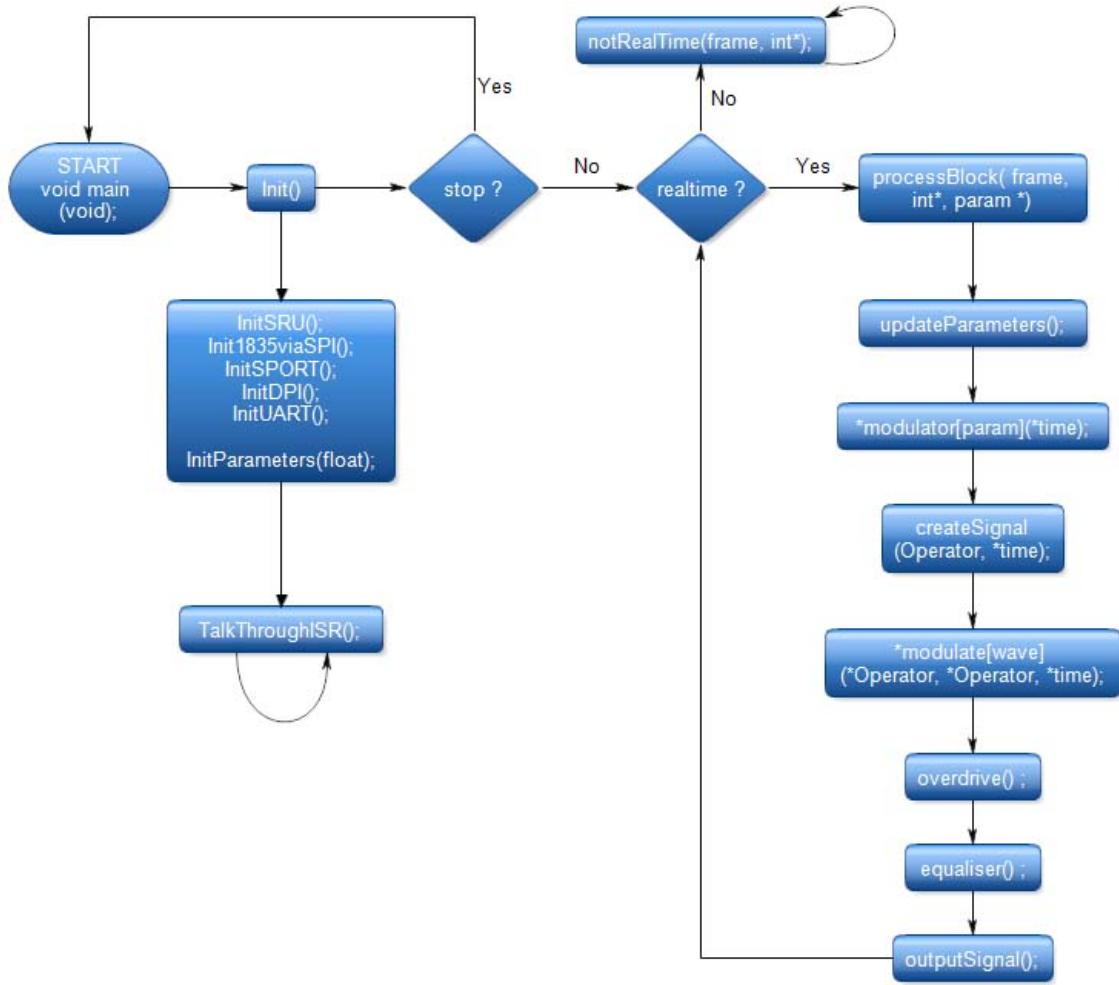


FIGURE 6.7: Flow chart of the DSP Synthesizer program
(Compare this figure with figure 2.12)

```

void notrealtime(unsigned int *block_ptr, int *time)
{
    int i;
    float scalar = powf(2.0,23.0);
    long j,N,readptr,ref;

    //Clear the Block Ready Semaphore, Set the isProcessing Semaphore before starting
    blockReady = 0;
    isProcessing = 1;

    for(i=0;i<NUM_SAMPLES;i++)
    {
        *(block_ptr+i) = 0.1*scalar*sin(2.0*3.141592*8000.0/sampleFreq * ( (int)(*time)++ ));
        (*time)%=sampleFreq;
    }

    //Clear the Processing Active Semaphore after processing is complete
    isProcessing = 0;
}

```

6.4.1.2 Realtime: let's play

In order to easily manipulate our sound processing algorithms we need the same arguments as indicated in section 6.4.1.1 plus an additional pointer used to communicate with the Labview interface.

Refresh parameters: Before we start to create our signal we first fetch the parameters as the user defines them in the Labview interface. Those values are put into global variables so they have a reserved spot in the program memory.

Select modulation type: Depending on which modulation algorithm the user selects a different function will be executed. This is almost literally the definition of function pointers, isn't it? Since the different modulation modes have a fixed value in the labview interface we can define the function pointers very efficiently. The definition of the function pointer we used is a function that takes one int pointer as argument and returns void. Below the creation and call of the (fancy) function pointers is shown:

```
Modulation *modulator[4] = {&TwoOperators, &ThreeOperators, &FourOperators,
&FourOperatorsPlus};

(*modulator[(int) parameters[0]])(time);
```

Create and modulate signal: For efficiency reasons we mostly pass arguments by reference. However in the case below we use pass by value because we don't want to adapt the value of the pointer immediately (and copying the value to a local variable is also a bit sloppy). So to create and modulate the signal first dereference the time pointer received by the function pointer:

```
void TwoOperators(int *t)
{
    int i;
    Operator *carrier = &op[0];
    Operator *modulator = &op[1];

    createSignal(modulator, *t);
    (*modulate[carrier->wave])(carrier, modulator, *t);
}
```

As you can see, the modulation algorithm also depends on which type of wave is selected and this is also fixed with function pointers. For the complete implementation of the modulate algorithms we refer to the file *operator.c*.

Overdrive: If the effect is enabled the function overdrive is called. This function takes as arguments operator 0 and a pointer to the effect's parameters. The creation and modulation algorithms are so that operator 0 always contains the whished signal, this is why it is used as argument here and in the next functions.

```
if( effect.enabled ) overdrive(&op[0], &effect);
```

Equaliser: If the equaliser is enabled the following function is called:

```
if( equaliser.enabled ) enableEqualiser(&op[0], &equaliser);
```

Output signal: Now we have fully modulated our signal we still have to put it into the TCBs to send it to DAC4. One last thing that is added here is a control of the overall gain of the signal.

We may not forget that because we wanted to use the timing values several times we used a pass by value. That's why after the whole frame is processed we still have to update the time reference.

```

    for(i=0; i<NUM_SAMPLES; i++)
    {
        *(block_ptr+i) = (unsigned int)(signed int)(op[0].signal[i] *
            Gain(&equaliser)*0.1*scalar );
    }
    (*time)+= 1024;
    (*time)%=sampleFreq;
}

```

6.4.2 Key functions

```

void createSignal(Operator *, int) {
    int i;
    switch(op->wave)
    {
        case SINE:
            for(i=0; i<NUM_SAMPLES; i++)
            {
                op->signal[i] = envelope(op) * sin( 2.0 * 3.141592 *
                    op->frequency / sampleFreq * t++ );
            }
            break;
        case TRIANGULAR:
            for(i=0; i<NUM_SAMPLES; i++)
            {
                op->signal[i] = 1-(fabs((t++*op->frequency%sampleFreq)-sampleFreq/2)
                    / (sampleFreq/4) - 1 );
                op->signal[i] *= envelope(op);
            }
            break;
        case SAWTOOTH:
            for(i=0; i<NUM_SAMPLES; i++)
            {
                op->signal[i] = 2*( t++*(op->frequency)%sampleFreq )/sampleFreq - 1;
                op->signal[i] *= envelope(op);
            }
            break;
        default:
            break;
    }
}

```

Remark: Why is there no $t\%=\text{sampleFreq}$ in this function?

void ModulateSine(Operator *carrier, Operator *modulator, int t) Please remember formula 2.10 discussed in chapter 2 on frequency modulation.

```

int i;
for(i=0; i<NUM_SAMPLES; i++)
{
    carrier->signal[i] = sin( 2.0 * 3.141592 * carrier->frequency /
        sampleFreq * t++ + (modulator->signal[i]) );
    carrier->signal[i] *= envelope(carrier);
}

```

Chapter 7

Communication with LabVIEW User Interface using UART Port Controller

Written by Thuy Pham

7.1 Introduction

Communicating with peripheral devices plays a vital role in system design in general and in DSP systems specifically. For this purpose, SHARC Processor uses Digital Peripheral Interface (DPI). The interface provides both connections to two serial peripheral interface ports (SPI) and two universal asynchronous receiver-transmitters (UARTs). In this report, Universal asynchronous receiver transmitters (UARTs) will be examined in relationship with LabVIEW User Interface. It is necessary to notice that the processor provide a full duplex universal asynchronous receiver/transmitter (UART) port, which is fully compatible with PC-standard UARTs. The UART port provides a simplified UART interface to other peripherals or hosts, supporting full duplex, DMA supported asynchronous transfers of serial data. The UART also has multiprocessor communication capability using 9-bit address detection. This allows it to be used in multi-drop networks through the RS-485 data interface standard. The UART port also includes support for five data bits to eight data bits, one stop bit or two stop bits, and none, even, or odd parity. More specifically, the UART port supports two modes of operation as mentioned below:

- **PIO (programmed I/O):** The processor sends or receives data by writing or reading I/O mapped UART registers. The data is double-buffered on both transmit and receive.
- **DMA (direct memory access):** The DMA controller transfers both transmit and receive data. This reduces the number and frequency of interrupts required to transfer data to and from memory. The UART has two dedicated DMA channels, one for transmit and one for receive. This is the mode that will be used for transferring data between LabVIEW program and DSP board.

7.2 UART data frame

When two UARTs communicate, both transmitter and receiver need to know the signaling speed. The receiver does not know when a packet will be sent (no receiver clock); thus, the protocol

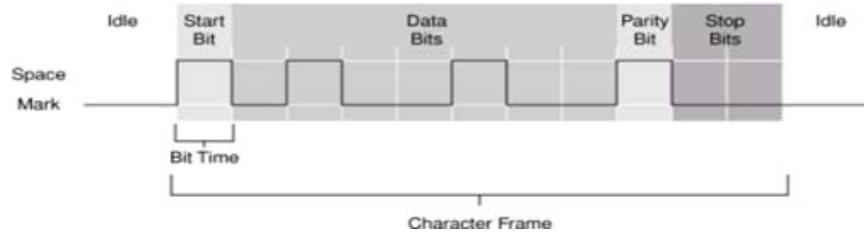


FIGURE 7.1: A typical UART data frame

is termed "asynchronous." This is very important because if we set up parameters (baud rate, start bit, stop bits, and parity bit) in both transmitter and receiver is incompatible, we will get incorrect data. In addition, the receiver circuitry is correspondingly more complex than that of the transmitter. The transmitter simply has to output a frame of data at a defined bit rate. Contrastingly, the receiver has to recognize the start of the frame to synchronize itself, and therefore determine the best data sampling point for the bit stream.

The processor supports a set of parameters' values for UART communication as below:

SST	Parameters	Value
1	Baud rate	2400, 4800, 9600, 19200, 38400, 57600, 115200, 921600, 6250000
2	Data bits	5 to 8
3	Stop bits	1 or 2
4	Parity	None, even, odd

7.3 UART external interface

The DSP processor communicates with peripheral devices in DSP board is described as the block diagram in figure 7.2:

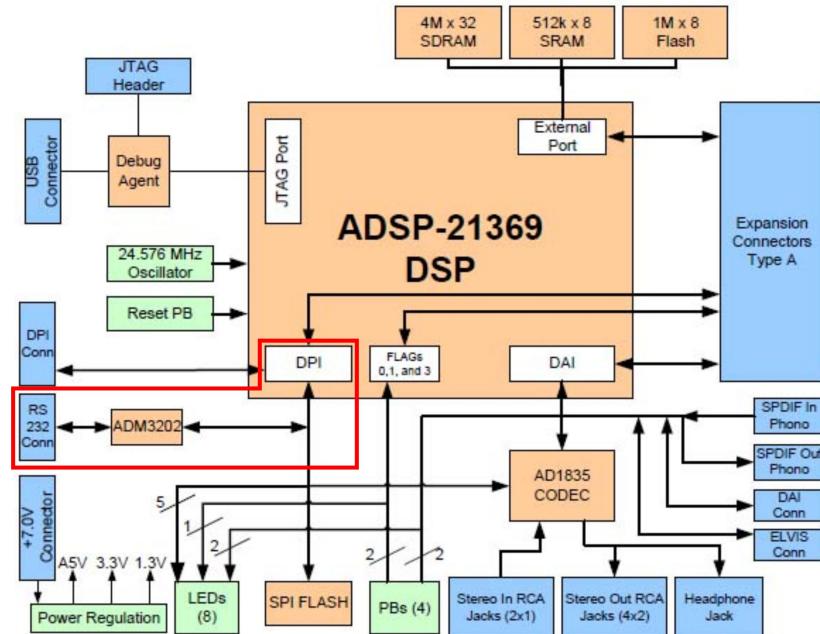


FIGURE 7.2: System Architecture Block Diagram

In this board, the EIA-232E interface is used to communicate with other external serial com-

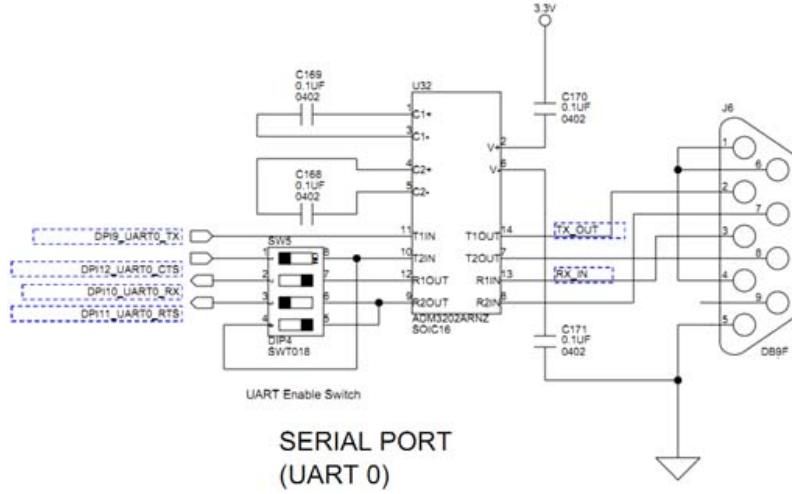


FIGURE 7.3: Schematic for connection of UART port

munication devices. Instead of using traditional circuit for compatible purpose of power levels, in the DSP board, manufacture uses ADM3202 chip. The significant features of the chip are low power consumption and can operate at data rates up to 460 kbps make them ideal for battery powered portable instruments and high speed requirement.

7.4 UART configuration in DSP Processor for communication

The DSP Processor provides a set of PC style, industry standard control and status registers for the UART.

Register Overview for UART Module:

- Line Control Register (UARTxLCR). Controls format of the data character frames. It selects word length, number of stop bits and parity.
- Divisor Latch High/Low Register (UARTxDLL, UARTxDLH). Characterize the UART bit rate. The divisor is split into the divisor latch low byte (UARTxDLL) and the divisor latch high byte (UARTxDLH).
- Mode Control Register (UARTxMODE). Controls packing and address modes.
- Transmit Buffer Control Register (UARTxTXCTL). Controls core or DMA operation.
- Receive Buffer Control Register (UARTxRXCTL). Controls core or DMA operation.
- Interrupt Enable Control Register. Enables interrupt requests from system handling.
- Line Status Register (UARTxLSR). Returns status of controls format of the data character frames as overrun or framing errors and break interrupts.
- Transmit Status Register (UARTxTXSTAT). Returns status of core or DMA operations.
- Receive Status Register (UARTxRXSTAT). Returns status of core or DMA operations.
- Interrupt Identification Status Register (UARTxIIR). The register is used to get the status of all interrupts into one channel.

To configure parameters for transmitting data between PC and the DSP board, the number of simple steps can be followed:

1. Route the UART to the DPI of DSP Processor

The routing is implemented by software, in particular in the VisualDSP++, it is in the SRU macro. It is necessary to notice that, data is transmitted and received by the least significant bit (LSB) first (bit 0) followed by the most significant bits (MSBs).

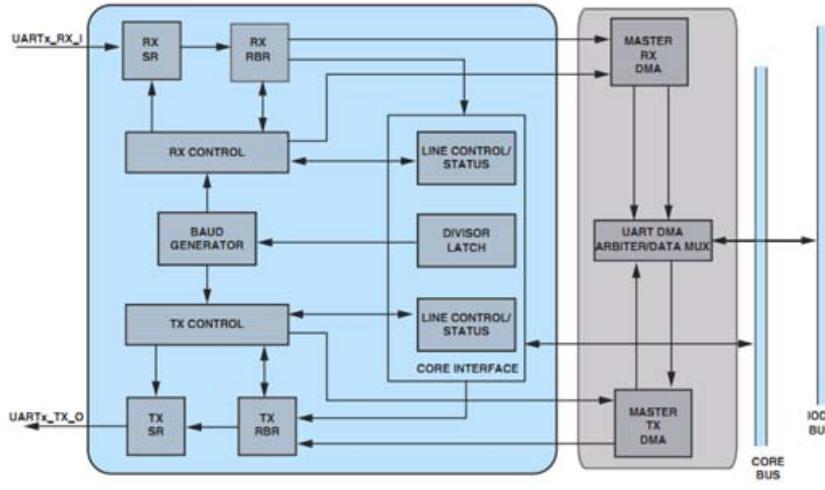


FIGURE 7.4: UART Functional Block Diagram

2. Set up baud rate, data bits, stop bits, parity bit

The bit rate is characterized by the peripheral clock (PCLK) and the 16-bit divisor. The divisor is split into the UART divisor latch low byte register (UARTxDLL) and the UART divisor latch high byte register (UARTxDLH). $\text{UART Baud rate} = \text{PCLK}/(16 \text{ divisor})$, where PCLK is the system clock frequency.

All data words require a start bit and at least one stop bit. With the optional parity bit, this creates a 7 to 12-bit range for each word. The format of received and transmitted character frames is controlled by the line control register (UARTxLCR).

3. Set up core mode of operation

This requires software management of the data flow using either interrupts or polling. Core transfers move data to and from the UART by the processor core. To transmit a character, load it into the UARTxTHR register. Received data can be read from the UARTxRBR register. The processor must write and read one character at time. To prevent any loss of data and misalignments of the serial data stream, the UART line status register (UARTxLSR) provides two status flags for handshaking - UARTTHRE and UARTDR. Core transfers through the UART is started by setting up and writing to transmit and receive control registers, enabling the module using the UARTEN bits in the UARTxTXCTL and UARTxRXCTL registers.

4. Set up the UART interrupt

The UART receive and transmit interrupts are programmed through the peripheral interrupt control registers (PICRx) as separate interrupts. (By default, these interrupts are not configured in the IRPTL register - the PICRx register has to be programmed to configure them.)

The UART interrupt enable register (UARTxIER) is used to enable requests for core system handling of empty or full states of UART data registers. When polling is used as a means of action, the UARTRBFIE and/or UARTTBEIE bits in this register are normally set.

7.5 Programming for communication between DSP Processor and LabVIEW User Interface

7.5.1 LabVIEW program for writing data

LabVIEW program provides a set of powerful tools for programming User Interface. Based on these things, user can make a program that clearly and visually shows parameters, graphs and results. LabVIEW also has many instrument drivers, which are a set of modular software functions that use the instrument commands or protocol to perform common operations with the instrument, for a variety of programmable instruments that use the GPIB, VXI, PXI, or serial interfaces.

With serial communication, in particular in RS-232, which is a standard developed by the Electronic Industries Association (EIA), LabVIEW provides some VISA functions for the target. They are located in **All Functions >> Instrument >> I/O >> Serial**. The VISA Configure

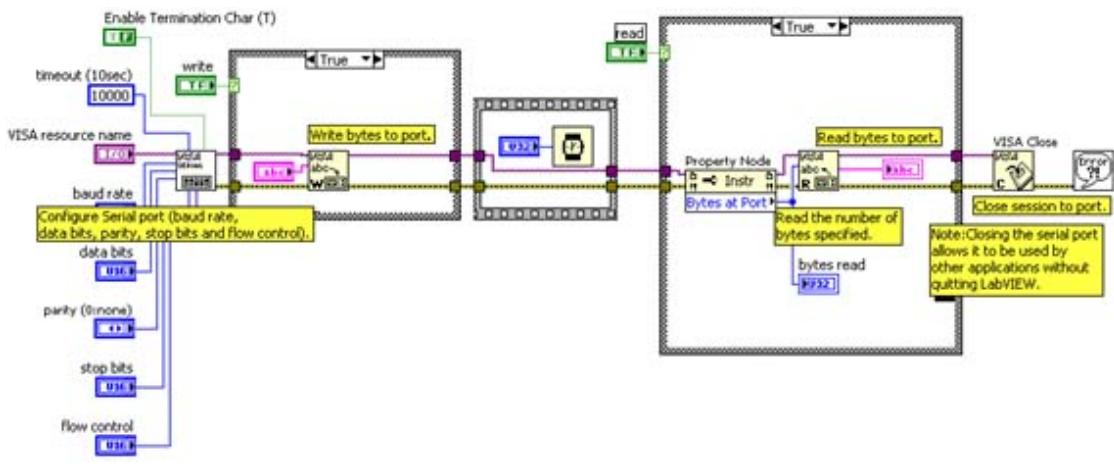


FIGURE 7.5: Write/Read a string to and from a COM port

Serial Port VI initializes the port identified by VISA resource name to the specified settings: timeout sets the timeout value for the serial communication; baud rate, data bits, parity, and flow control specify those specific serial port parameters.

1. The VISA Configure Serial Port VI initializes the port identified by VISA resource name to the specified settings: timeout sets the timeout value for the serial communication; baud rate, data bits, parity, and flow control specify those specific serial port parameters.
2. The VISA Write function sends the string.
3. The VISA Read function reads back up to bytes into the buffer, and the Simple Error Handler VI checks the error condition.
4. The VISA Close function terminates the communication channel to the instrument and deal locates the resources for the DSP.

Depending on the purpose, the string can be written to the device in different forms. As the subVI described below, the string is combined by three parts, which are identified parameter order (two characters), the value needed to pass, and the character that uses for identifying read or write (character "w"). It is important to notice that maximum character transmission rate is depended on the baud rate and the bits per frame. More specifically, this rate is just the baud rate divided by the bits per frame. So, when there is a consecutive transfer required, the number of characters per string needs to be cared. Otherwise, there will be data lost. Based on the prime

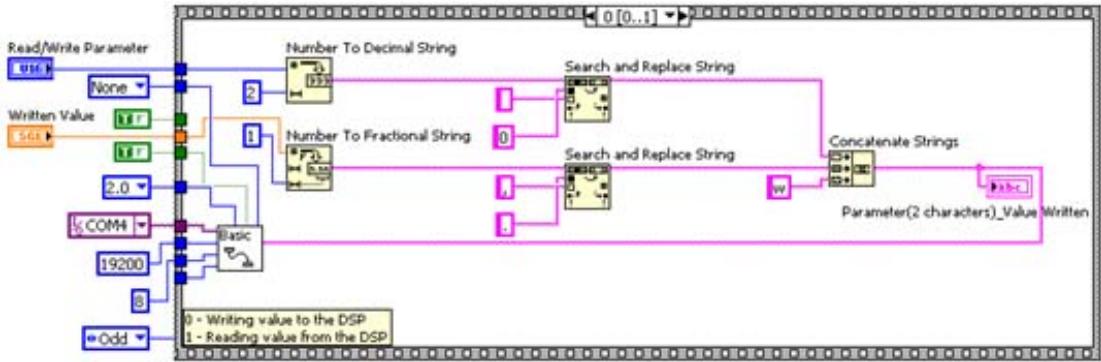


FIGURE 7.6: Formatting a string before writing to DSP

subVI as showed in figure 7.6, there are a number of ways to program an application that writes and reads multiple parameters to and from a DSP Processor. It needs to be reminded that, the principle of transferring data between DSP and PC is based on the interrupt. In other words, every time data is written to the DSP from PC, the serial DSP interrupt occurs and then its ISR (interrupt service routine) will run and process the received data. As a result, it is impossible to write more than one parameter at the same time that the DSP can process independently. Of course, each parameter has a different purpose. Because of this, the parameters should be organized so that only one parameter can write its data to the DSP at the time. It is very important for copying data between parameters that will be studied later.

Figure 7.7 shows an example program that wants to write multiple parameters to the DSP.

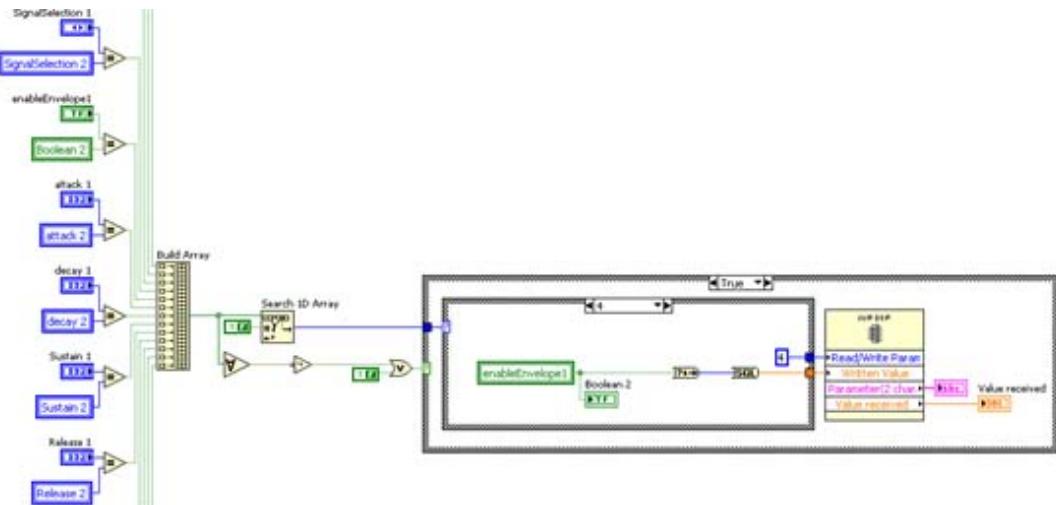


FIGURE 7.7: A part of writing parameters to Operator 1 program

The idea of the program is that every time a parameter changes its value, which is detected due to search 1D array, the read/write part starts working. Otherwise, the program does nothing. This reduces a lot of works for the DSP Processor. Then, the latest value of the parameter and its order will be combined into string before writing to the processor. Notice that, LabVIEW also provides some data convert functions so that different data types of parameters can work properly.

7.5.2 UART Interrupt

There are three things needed to be cared while using UART interrupts. First, you must map the UART interrupts to one of the interrupts in the interrupt vector table. In particular, the UART has to be mapped to exactly the peripheral interrupt sources. This can be achieved by changing the default source of the peripheral interrupt priority control register with the UART source, using the interrupt select values of the UART receive interrupt. The receive interrupt select values for UART0 (using in our application program) and UART1 are 0x13 and 0x14, respectively.

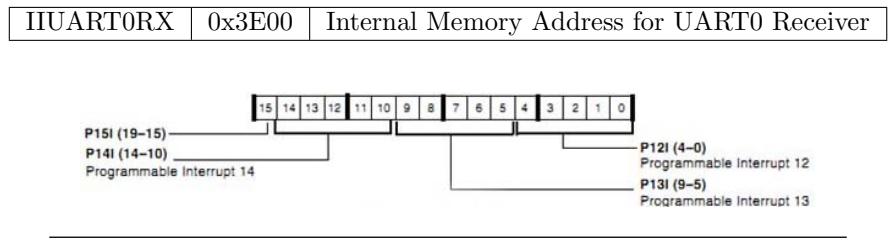


FIGURE 7.8: PICR2 Register

```
*pPICR2 &= ~(0x3E0); // Sets internal memory address for UART0 receiver
*pPICR2 |= (0x13<<5); // Sets default value (0x13) for the URATO receive interrupt
```

The second thing is that enabling the UART interrupts internally by setting the corresponding bits in the UART interrupt enable register (UARTxIER). This needs to be done after all the UART settings (such as word length, parity, and so on) have been programmed, because in transmit mode as soon as the transmit buffer empty bit is enabled in the UARTxIER register, it vectors to the interrupt. If this bit is enabled before any of the UART settings are programmed, the data transmitted in the transmit interrupt service routine does not comply with the UART settings that are programmed later, leading to a communication error.

```
*pUART0IER = UARTRBFIE; // Enables UART0 receive interrupt
```

The third, which is also the most important, is that using C Interrupt Handler with own Interrupt Service Routine. C provides its own set of interrupt handlers via the interrupt() and signal() functions. VisualDSP++ has extended the interrupt() function with **interruptf()** and **interrupts()** versions. The interrupt handler consists of three parts - the initialization, the interrupt vector, and the interrupt service routine.

```
interrupt(SIG_P13, UARTisr);
```

Initialization performs the tasks previously associated with the C interrupt() function - assigning the interrupt service routine to an interrupt vector (dynamic ISR vectors only), unmasking the interrupt, and enabling global interrupts. The interrupt vector routes execution to the appropriate interrupt service routine. This routine must be compatible with the procedure of writing data of LabVIEW program. Otherwise, the data that DSP Processor expected to receive is different with data transmitted.

```
// Read the data from the buffer
// Receive Buffer Registers (UARTxRBR)
value = *pUART0RBR;

/* Echoes back the value on to the hyperterminal*/
```

```

// Wait until it is sent
while ((*pUARTOLSR & UARTTHRE) == 0){ ; }

// Writting a string to DSP
// The character in buffer is "w" (ASC-II of "w" = 0x77 equivalent to 119 decimal)
if(value == 119)
{
    for(i = 0; i < 2; i++) address[i] = inputstream[i];

    // The two first characters writting from LabView is the index
    // int atoi(const char *str); - Convert string to integer
    index = atoi(address);

    // The rest of the writting string is value (with a character - w)
    // Double atof ( const char * str ); - Convert string to double
    // 2 - is the place of starting value string
    writtenValue = atof(inputstream + 2);

    // Writting value to the global variable
    algoparam[index] = writtenValue;

    // Adds data to the buffer
    // Transmit Holding Registers (UARTxTHR)
    *pUARTOTHR = value;

    teller = 0;
}

```

7.5.3 The FM synthesizer application

The application is programmed to provide two services. First, the program can write the values of parameters to the DSP so that DSP board can produce different sounds. Secondly, the LabVIEW User Interface also can show a number of features of the signal, such as the form of the envelope, frequency domain and the spectrogram.

To provide the features above, the application is programmed following a hierarchy as figure 7.9. There are four main parts programmed in the application, which are programs for different

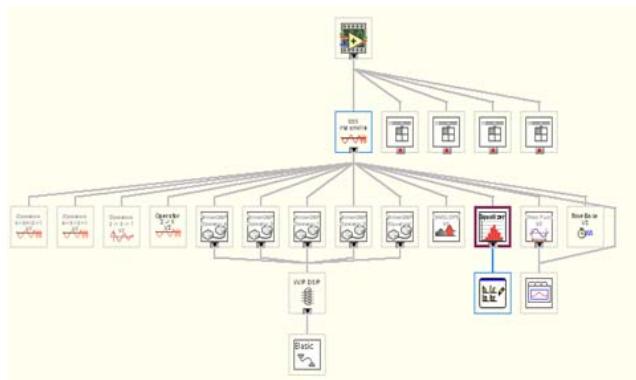


FIGURE 7.9: VI Hierarchy of LabVIEW User Interface

algorithms of combining operators, programs for the envelope, programs for equalizer, sound effect and display, and writing data to the DSP board. The first part includes four algorithms: two operators, three operators, four operators modulated straightforward and four operators

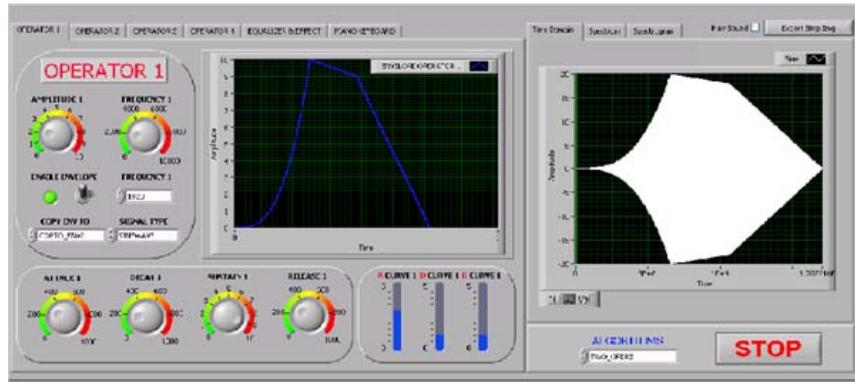


FIGURE 7.10: Front Panel of the application

that actually consist of a pair of two operators. This part can be expanded in future to have more possibilities for producing sounds. The second part of the application is the envelope.

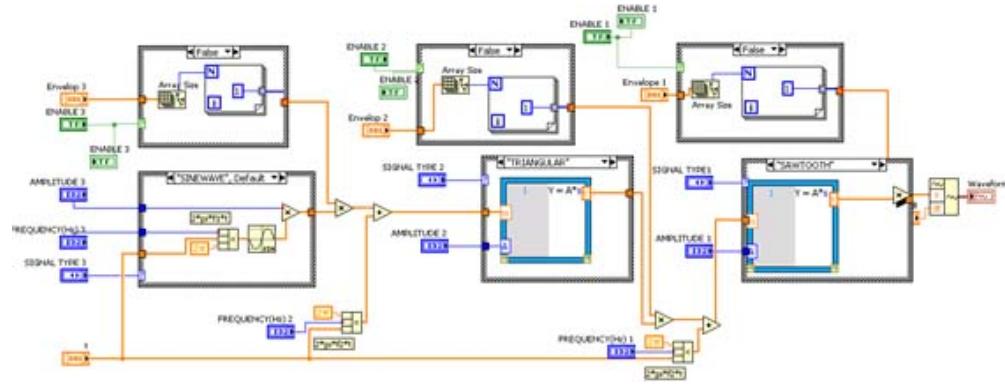


FIGURE 7.11: The algorithm for three operators

It is used to modify the shape of the signal so that the output can have different frequency components. The third is programmed to produce more flexibility for the application. With this, user can choose specific range of frequency of the output signal (using the equalizer) and also can make distortion (using the effect part). The last one is writing data part, which is discussed in "LabVIEW program for writing data", see section 7.5.1 above. More detail about the application, subVIs is available in appendix part in this report.

7.6 Conclusion

This report discusses a number of things, which are needed to take into account when working with UART controller of the DSP processor and LabVIEW program. From my point of view, despite the fact that each part plays a different role to make the application work properly, the most important thing is that the data (a string of characters) need to be define completely the same in LabVIEW User Interface and UART interrupt service routine. Otherwise, the data that DSP Processor wants to receive for processing is not expected.

To improve quality of communication between LabVIEW User Interface and DSP board, context switch of UART interrupt needed to be considered. However, because of time and the limited ability, it could be studied in the future.

Chapter 8

Critical reflections

Written by Bjorn Deraeve

8.1 Implementation problems

In chapter 6 we already explained how the program on the SHARC processor is designed. Here we will have a closer look at some problems encountered during the development.

8.1.1 Communication with the Labview interface

As explained in chapter 7, the parameters read from the interface are interrupt driven. As a consequence the program on the SHARC processor does not receive values if the user doesn't change any particular input on the interface.

For normal use this is not a problem since the values are initialized correctly when the program is started. However during the debugging process this often led to confusion and searching on problems in the code that actually weren't there.

8.1.2 Combining different program parts and general programming issues

Including smaller topics like the envelopes, effects and equaliser into the main program sometimes led to major problems. Next to the inevitable minor adjustments in the program parts also the interaction with the interface caused problems. However, most of these problems were small mistakes that could easily be fixed. Nevertheless not all team members succeeded into fixing those little errors.

Sometimes it was necessary to change some details of the main program in order for the sub-programs to work fluently. However, in order to do so efficiently one must at least understand the template we received at the beginning of the project and think well before changing stuff. There had to be intervened several times because a team member unknowingly stopped the whole program instead of just disabling his own algorithm. For example, as explained in chapter 6, processors on embedded systems need a `while(TRUE)` loop. Another useful thing to realize is that the periodic functions generating the waveforms need a variable that keeps track of the time. Updating this time in a local variable does not have the same effect like updating the global variable and again stops the program!

Also the modulo %-operator is common in DSP algorithms and all team members should its meaning. Looking this up with google would have been a great idea!

Finally before asking to include a subprogram into the main program the smaller part must first have been debugged for syntax errors (since notepad doesn't highlight such errors).

Next, some other concrete programming issues are described:

Sampling frequency: The template we started from mentionned a sampling frequency of 48 kHz. However that project used a stereo audio channel which meant the samples were sent to the AD1835A at a rate of 96 kHz. Because we filled the algorithm's frames with mono channel samples (instead of two samples per calculation) our sampling frequency in software matched the sampling frequency of the hardware.

Modulation: Enabling modulation caused some minor timing problems, however this could simply be fixed by passing the function's argument by value instead of by reference.

Envelopes: There were several problems with the envelope. In a first stage nothing worked and a trick must be implemented to create envelopes with a duration longer than one second. To do so the function received a counter in its arguments.

The updating of this counter was at first done in the main synthesize sequence, when the program's time was an integer multiple of the sampling frequency. However because of the fixed frame length of 1024 samples this happen only every 4 seconds. To avoid this the updating of counter was originally moved to the create signal function. This has no influence and even made things worse actually. The problem with placing the flag here was that if modulation is enabled the flag updated several times. However this whole time the counter argument worked and the envelope refused working even without modulation.

The underlying cause was that the envelope function made use of the wrong envelopes in memory so the parameters had no influence. Unfortunately not all team members were able to find such mistakes

Finally to fix all little problems the envelope received its own space in memory for keeping track of time. So now this isn't done anymore in the synthesize function but in the envelope itself. Because now the time is obviously updated instantaneously all problems are fixed.

8.1.3 Connecting the DSP board and use of the IDEs

There were numerous of problems with Analog Devices' IDE. Personally I could not connect with the board on my normal Windows. Running the program from within a virtual machine on that same computer seemed to magically solve the problem (after a few years of internetting). If we have to invent an explanation I'd say that Sony tampered with the OS' USB drivers to make them more energy efficient and thus not supporting the DSP board.

Getting VisualDSP++ and Labview to run fluently also took a lot of time. For VisualDSP++ an activation code had to be registered. This became a problem since by the end of the project we got warnings that the activation would expire in less than a few days. Another issue with this activation code was that in week 10 a team member asked for it while it had already been e-mailed to all team members in week 4.

VisualDSP++ also seemed to be a very hard program to work with since a team member needed step-by-step click instructions for how to activate the editor's line numbers.

For Labview there were no valid licenses at all so the program had to be reinstalled a few times. Also to use all aspects of the interface several extra Labview plugins needed to be installed and again the only way to do is by completely reinstalling Labview, this became more or less a routine.

8.1.4 Extreme programming

Due to the tricky situation with the interface several problems asked a lot of time to get fixed or didn't get fixed at all. For those harder to find errors a second pair of eyes would have been welcome.

Though in an attempt to improve the overall impression of the project it was best to do have look at those extra problems and happily been able to fix some of them. *However it cannot be that there is only one person responsible to make all things work properly, the lack of a good programming partner was big. With such a partner a lot of improvements could have been made, some of them are discussed in section 8.4.*

8.1.5 Latex

Latex' standard color package does not support the color pink. My disappointment was huge, after heavy consideration I will declare this as the biggest lack encountered during this EE5-project.

8.2 Team problems

Next to common problems teams sometimes encounter like poor communication, group thinking and difficulty making decisions we also suffered from other problems. They are described shortly below:

- Major issues

- Lack of basic programming knowledge
- Update files to the shared folder
- Lack of inspiration for simple program features

- Minor issues

- Update timesheets in time
- Labview vs. C distinction
- Give feedback on the programs: this only happened for the interface
- Chaotic organisation of files in the shared folder

8.3 Achieved activities

Student	Hours
Bjorn Deraeve	210
Frederik De Greef	140
Kenneth Piot	140
Thuy Pham	180

TABLE 8.1: Project logbook

Student	Topic
Bjorn Deraeve	Admin: Meeting reports Labview: Connect the basic interface C: General program structure C: Signals and modulation C: Import and adapt (fix) the other programs C: Debug and fix envelope program Matlab: plots for report Labview: Bessel functions L ^A T _E X: cls, bib and tex template, chapters 1,2,5,6,8
Frederik De Greef	C & Matlab: effects C & Matlab: equaliser Report: chapter 4
Kenneth Piot	C: Envelopes Labview: Piano interface Report: chapter 2
Thuy Pham	Labview: full interface C: Configuration of UART communication Labview: waves and spectra Labview: envelopes and graphs Labview: equaliser Labview: piano interface Report: chapter 7

TABLE 8.2: Activities

8.4 Future work

This section summarizes some features that could add additional value to DSP Synthesizer.

Envelopes: It is no coincidence the piano interface is located in the chapter on the envelopes. Combining the piano keystroke with the envelopes should have been an easy to implement extension.

Equaliser: An extended equaliser would provoke the musicians' creativity more.

SDRAM: Being able to save created sounds and reuse them to modulate signals. Also prestored effects could be offered this way. Search how to initialize the chip, ...

SHARC-ADSP 21369 IRQs: Using the board's buttons to increase or decrease the volume. This is an interrupt driven event. Combining this interrupt with the other, time pressured interrupts, would have presented some interesting programming problems.

Hardware controller: Instead of using the labview interface use a real controller with a more advanced communication protocol. The board supports SPI, I²C,...

Chapter 9

Bibliography

- [1] Cristina Alcaraz, Pablo Najera, Javier Lopez, and Rodrigo Roman. *Wireless Sensor Networks and The Internet of Things: Do We Need a Complete Integration?* Computer Science Department University of Malaga, Spain, 2010.
- [2] Atmel. *8-bit Atmel Microcontroller with 64K-128K-256K Bytes In-System Programmable Flash*, October 2012.
- [3] Isidor Buchmann. Elevating self-discharge, 2013. http://batteryuniversity.com/learn/article/elevating_self_discharge/.
- [4] Sebastian Buttrich. Wireless sensor networks, 2010. Course Lecture SPVC2010.
- [5] Cooking Hacks by Libelium Comunicaciones Distribuidas S.L. Wireless sensor networks open source platform. <http://www.cooking-hacks.com/index.php/documentation/tutorials/waspmove/>.
- [6] Digi International Inc. *X-CTU Configuration & Test Utility Software*, 90001003a edition, August 2008.
- [7] Digi International Inc. *XBee XBee-PRO ZB RF Modules*, 90000976l edition, April 2012.
- [8] Robert Faludi. *Building Wireless Sensor Networks*. O'Reilly Media Inc., first edition, 2011.
- [9] Kenneth Flamm. *Creating the Computer*. The Brookings Institution, 1988.
- [10] Wireless Sensor Networks Research Group. Libelium launches new generation of waspmote sensor nodes. April 2013. <http://www.sensor-networks.org/index.php?page=1309909904/>.
- [11] Libelium. 50 sensor applications for a smarter world. get inspired! May 2012. http://www.libelium.com/50_sensor_applications/.
- [12] Libelium Comunicaciones Distribuidas S.L. *Agriculture 2.0 Technical Guide*, 2-0.3 edition, November 2012.
- [13] Libelium Comunicaciones Distribuidas S.L. *Gases 2.0 Technical Guide*, 2-0.2 edition, November 2012.
- [14] Libelium Comunicaciones Distribuidas S.L. *Waspmove Datasheet*, 2.3 edition, November 2012.
- [15] Libelium Comunicaciones Distribuidas S.L. *Waspmove Interruptions Programming Guide*, 0.2 edition, June 2012.
- [16] Libelium Comunicaciones Distribuidas S.L. *Waspmove Power Programming Guide*, 0.6 edition, July 2012.

- [17] Libelium Comunicaciones Distribuidas S.L. *Wasp mote RTC Programming Guide*, 0.4 edition, June 2012.
- [18] Libelium Comunicaciones Distribuidas S.L. *Wasp mote Technical Guide*, 3.4 edition, November 2012.
- [19] Libelium Comunicaciones Distribuidas S.L. *Wasp mote Utilities Programming Guide*, 0.2 edition, June 2012.
- [20] Libelium Comunicaciones Distribuidas S.L. *Wasp mote ZigBee Networking Guide*, 1.0 edition, November 2012.
- [21] Libelium Comunicaciones Distribuidas S.L. *Catalogue Wireless Sensor Networks 2013*, 4.2 edition, February 2013.
- [22] Libelium Comunicaciones Distribuidas S.L. *Sensor Compatibility*, 0.5 edition, January 2013.
- [23] John A. Stankovic. *Wireless Sensor Networks*. University of Virginia, Department of Computer Science, 2006.
- [24] Arduino Team. Xbee shield. <http://arduino.cc/en/Main/ArduinoXbeeShield/>.

Appendix A

XBee join and send time experiments

A.1 Distance relation test

This table contains the results of two different test programs (blue and green values). The program scenario is as follows:

1. Turn on Waspmove and XBee
2. Wait until XBee has joined
3. Measure battery level, temperature, humidity and pressure (each with 100 milliseconds delay between samples)
4. Send the values and turn off Waspmove

DISTANCE	ON	JOINED	MEASURED	SENT	JOIN TIME	SENDING TIME	
1 FLOOR	26	3082	7325	7461	3056	136	
	26	3082	7325	7481	3056	156	
	26	3080	7323	7760	3054	437	
	26	3082	7325	7467	3056	142	
	26	3082	7325	7467	3056	142	
	26	2462	6705	6738	2436	33	
	26	2462	6705	6742	2436	37	
	26	2460	6703	6864	2434	161	
	26	2462	6705	7677	2436	972	
	26	2462	6705	7279	2436	574	
	26	2462	6705	6866	2436	161	
							AVG SEND
							268

A.2 First time reduction

The next results are obtained via the following scenario:

1. Turn on Waspmot and XBee
2. Measure battery level, temperature, humidity and pressure (each with 100 milliseconds delay between samples)
3. Check XBee association
4. Send the values and turn off Waspmot

It appears that when the XBee has been on for a sufficient amount of time, the time to request the node's association state is constant at about 450ms.

A second conclusion is that without obstacles and with a node that already is joined a while before trying to send leads to a constant sending time.

ON	MEASURED	JOINED	SENT / HIBER	JOIN TIME	SENDING TIME	
26	4742	5194	5309	452	115	
26	4742	5194	5307	452	113	
26	4742	5192	5305	450	113	
26	4742	5194	5307	452	113	
26	4742	5194	5307	452	113	
26	4742	5194	5307	452	113	
26	4742	5194	5307	452	113	
26	4742	5194	5309	452	115	
26	4742	5194	5303	452	109	
26	4742	5194	5305	452	111	
26	4742	5194	5266	452	72	
26	4742	5194	5309	452	115	
26	4742	5194	5677	452	483	
26	4742	5194	5788	452	594	
26	4742	5194	5437	452	243	
26	4742	5194	5307	452	113	
26	4742	5194	5307	452	113	
26	4742	5194	5307	452	113	
26	4742	5194	5266	452	72	
26	4742	5194	5309	452	115	
26	4742	5194	5309	452	115	AVG SEND
						158

FIGURE A.2: XBee join and send times: first time reduction